

Computer Architecture I

CS-GY-6133

Topic: Introduction and Instruction Set Architectures (ISA)

Instructor: Siddharth Garg

Logistics

- Lectures
 - Fridays, 6PM-8:30PM, Pfizer Auditorium, Dibner
 - Quizzes/exercises to keep you involved (bring your laptops!)
 - 15 minute break to keep you refreshed
 - **Tweet your questions @csgy6133 during class**
- Office hours
 - Fridays 1PM-3PM
 - 2 Metrotech, 10.076 or **Skype ID siddharth.j.garg**
- TAs
 - Team of 6 TAs. Names and office hours will be announced soon.

Grading (HWs and Labs)

- 3 Homeworks (20% of your final grade)
 - First HW to be released 9/16, due 9/30
 - Groups of two
- 3 programming labs (20% of your final grade)
 - First lab. to be released 9/30; due in 2 weeks
 - Groups of two (same as HW groups)
- Important: submit your partner's name before next class on NYUClasses (under "Assignments")

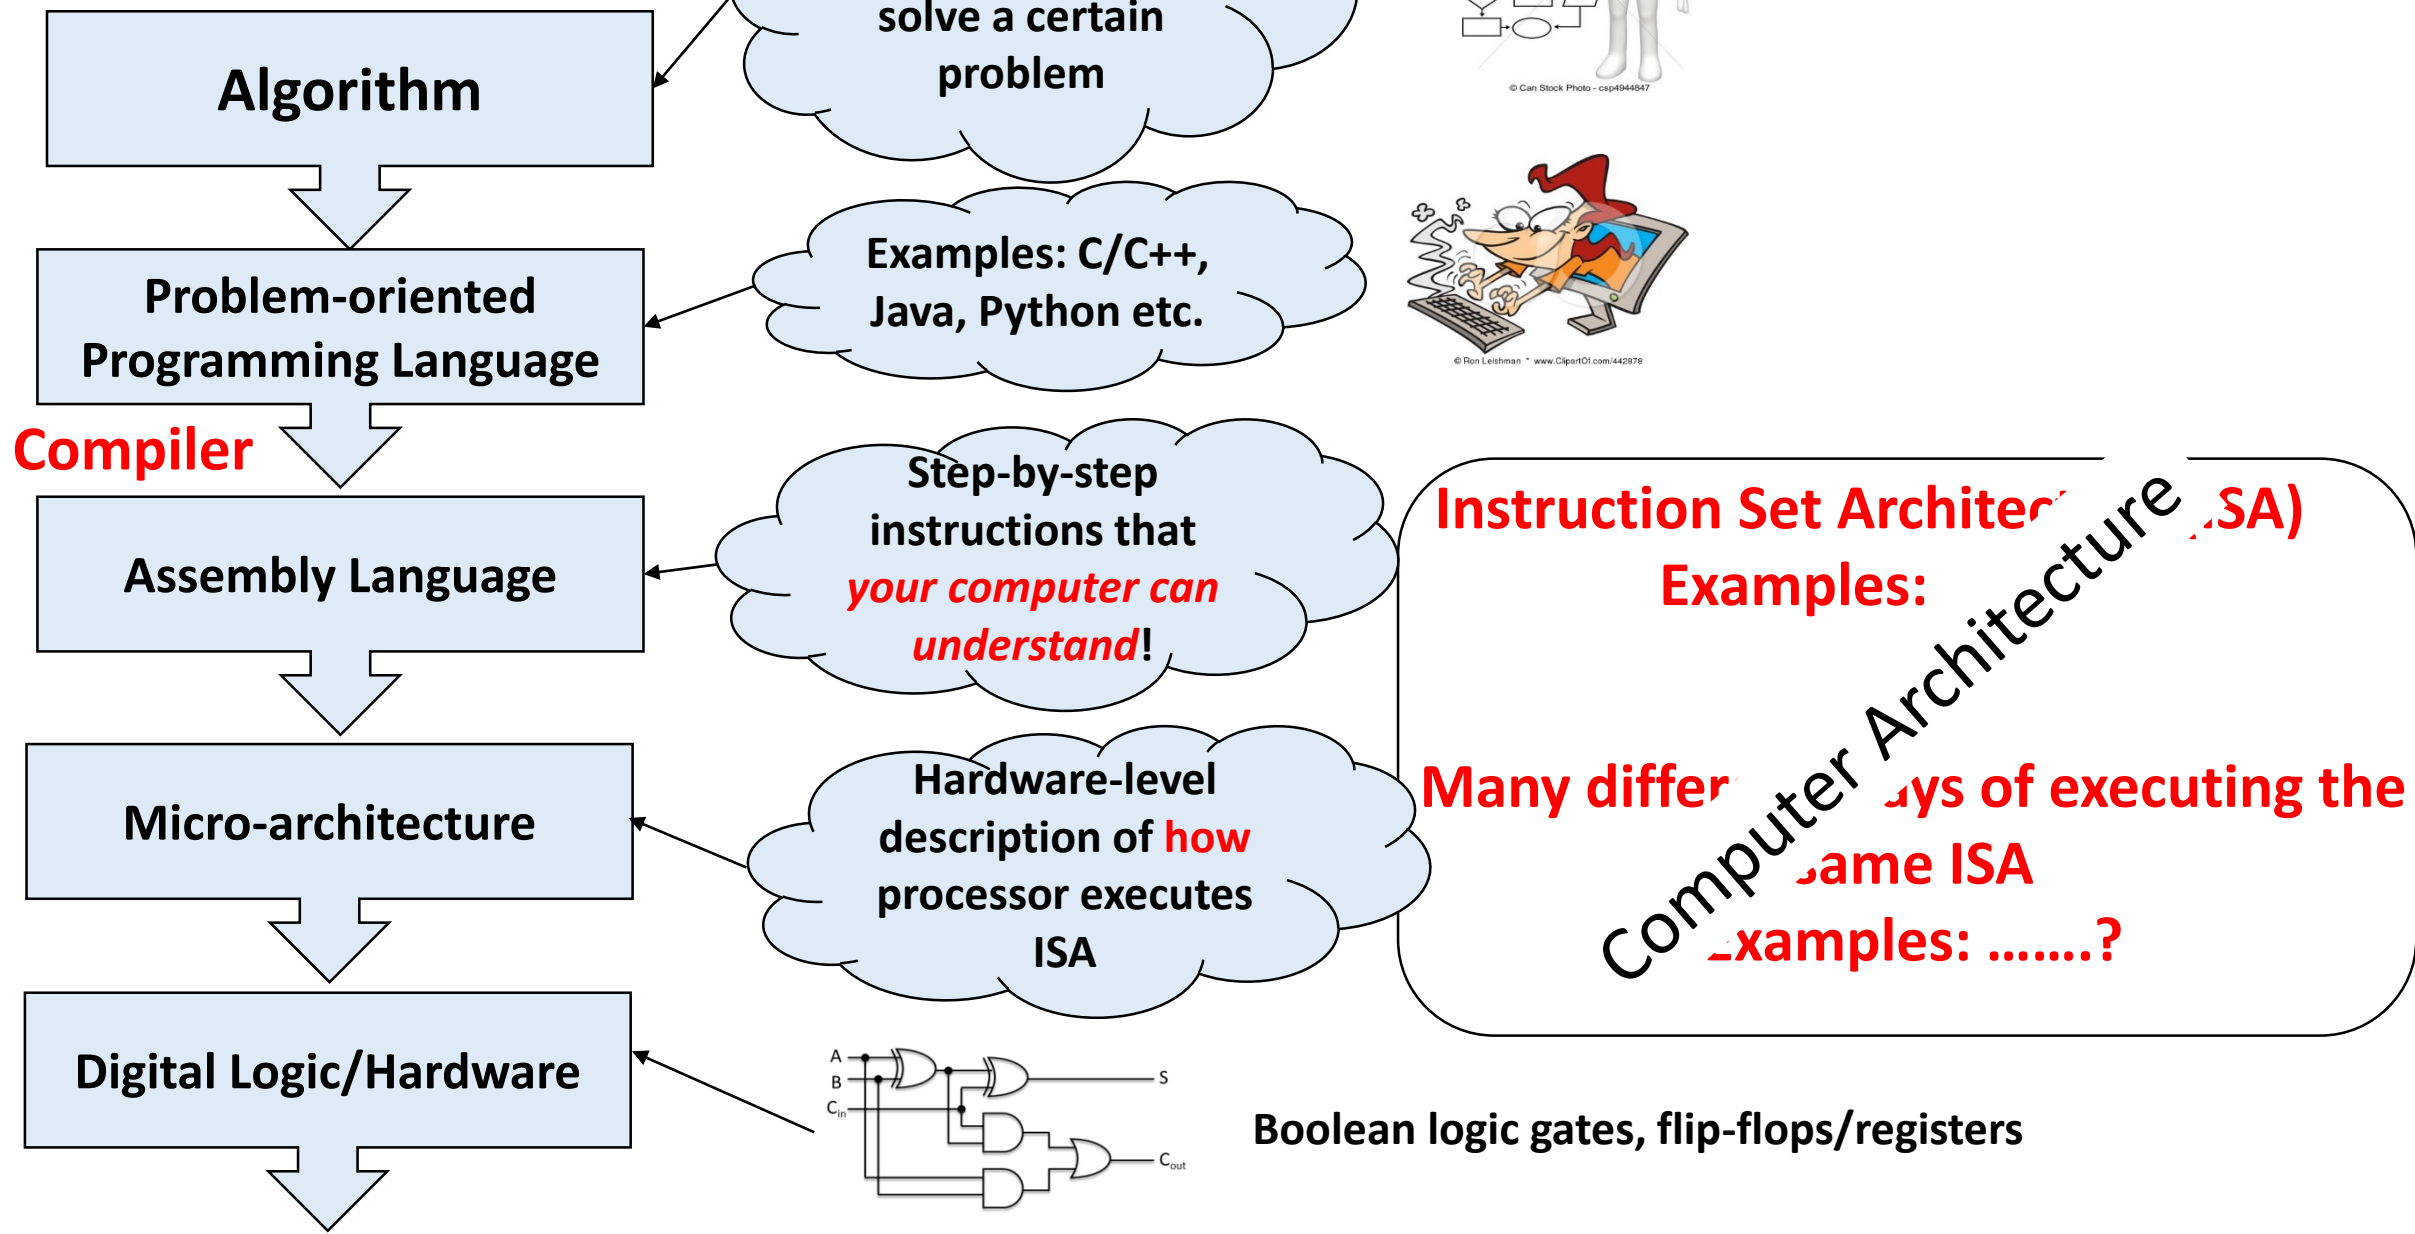
Grading (Midterm and Final)

- Midterm scheduled for Nov. 11th during class hours (25% of total grade)
 - Location will be announced later
- Final exam scheduled for Dec. 16th during class hours (35% of total grade)
 - Note: officially, final exam is scheduled for Dec. 23rd but kills your (and my) Christmas break
 - Location will be announced later

Plagiarism


- Don't do it
- Don't do it
- Don't do it
- Don't do it
- Don't do it
- Don't do it
- Don't do it
- Don't do it
- Why not?
 - You won't learn much
 - You'll likely get caught
- What happens if I get caught?
 - Lose 50% of your total course grade on first strike
 - Second strike and you fail the class

Introduction



Historical Perspective: Harvard Mark I

- 60 “constant registers” (23 decimals + sign)
 - Manually programmed constants
- 72 “storage counters” (23 decimals + sign)
 - To hold intermediate results of computation



- Instruction Format
 - 24 decimals
 - Fed with punch card

OUT (operand)	IN (operand, result stored here)	MISC. (which operation is performed)
Storage Counter 1 (Decimal Code: 1)	Storage Counter 2 (Decimal Code: 2)	SUB (Decimal Code: 32)

$SC2 \leftarrow SC2 - SC1$
No branch instructions!

- “Harvard Architecture” - separate memory for instructions and data

EDVAC: Electronic Discrete Variable Automatic Computer


- Internal magnetic “tape” based memory consisting of 1024 44-bit words
 - **Digital** (0,1) instead of decimal (0,1,2,..9) representation
 - Each word could represent either data or an instruction
 - Up to 16 different instructions, including **branch instructions**

Op Code (4 bits)	OP 1 (10 bits)	OP 2 (10 bits)	OP 3 (10 bits)	Next Instruction
Which operation (add, sub, mult, branch...)	Address of 1 st operand	Address of 2 nd operands	Address where result stored	Address from where next instruction fetched

- **Von Neumann Architecture**: Unified memory for data and instructions


Harvard Vs. Von Neumann Architecture

Von Neumann



- Memory must be a RAM (random access memory) with read/write capability

Harvard



- Data memory must be a RAM (random access memory) with read/write capability
- Instruction memory can be a ROM (read only memory)

Harvard Vs. Von Neumann: Pros/Cons

- Harvard Architecture

- **Con**: dedicated memory for code results in inefficient use of memory capacity
- **Pro**: Code memory can be read only (faster, smaller, cheaper)
- **Pro**: code is protected

- Von Neumann Architecture

- **Pro**: flexible use of available memory capacity
- **Con**: Memory needs to support read/write
- **Con**: code can be maliciously re-written

“Lecture 1: In-class Problem 1” on newclasses.nyu.edu under “Assignments”

Invention of the Transistor

- Vacuum tubes ruled in first half of 20th century Large, expensive, power-hungry, unreliable
- 1947: first point contact transistor
 - John Bardeen and Walter Brattain at Bell Labs




ENIAC had 18000 vacuum tubes and consumed 150 kW of power




From Transistors to Integrated Circuits

- 1958: First integrated circuit
 - Flip-flop using two transistors
 - Built by Jack Kilby at Texas Instruments
- 2010
 - Intel Core i7 μ processor
 - 2.3 billion transistors
 - 64 Gb Flash memory
 - > 16 billion transistors



Courtesy Texas Instruments




[Trinh09]

© 2009 IEEE

Moore's Law: Then

- 1965: Gordon Moore plotted transistor on each chip
 - Fit straight line on semilog scale
 - Transistor counts have doubled every 26 months



[Moore65]
Electronics Magazine

Integration Levels


SSI: 10 gates

MSI: 1000 gates

LSI: 10,000 gates


VLSI: > 10k gates

And Now...




Feature Size

- Minimum feature size shrinking 30% every 2-3 years

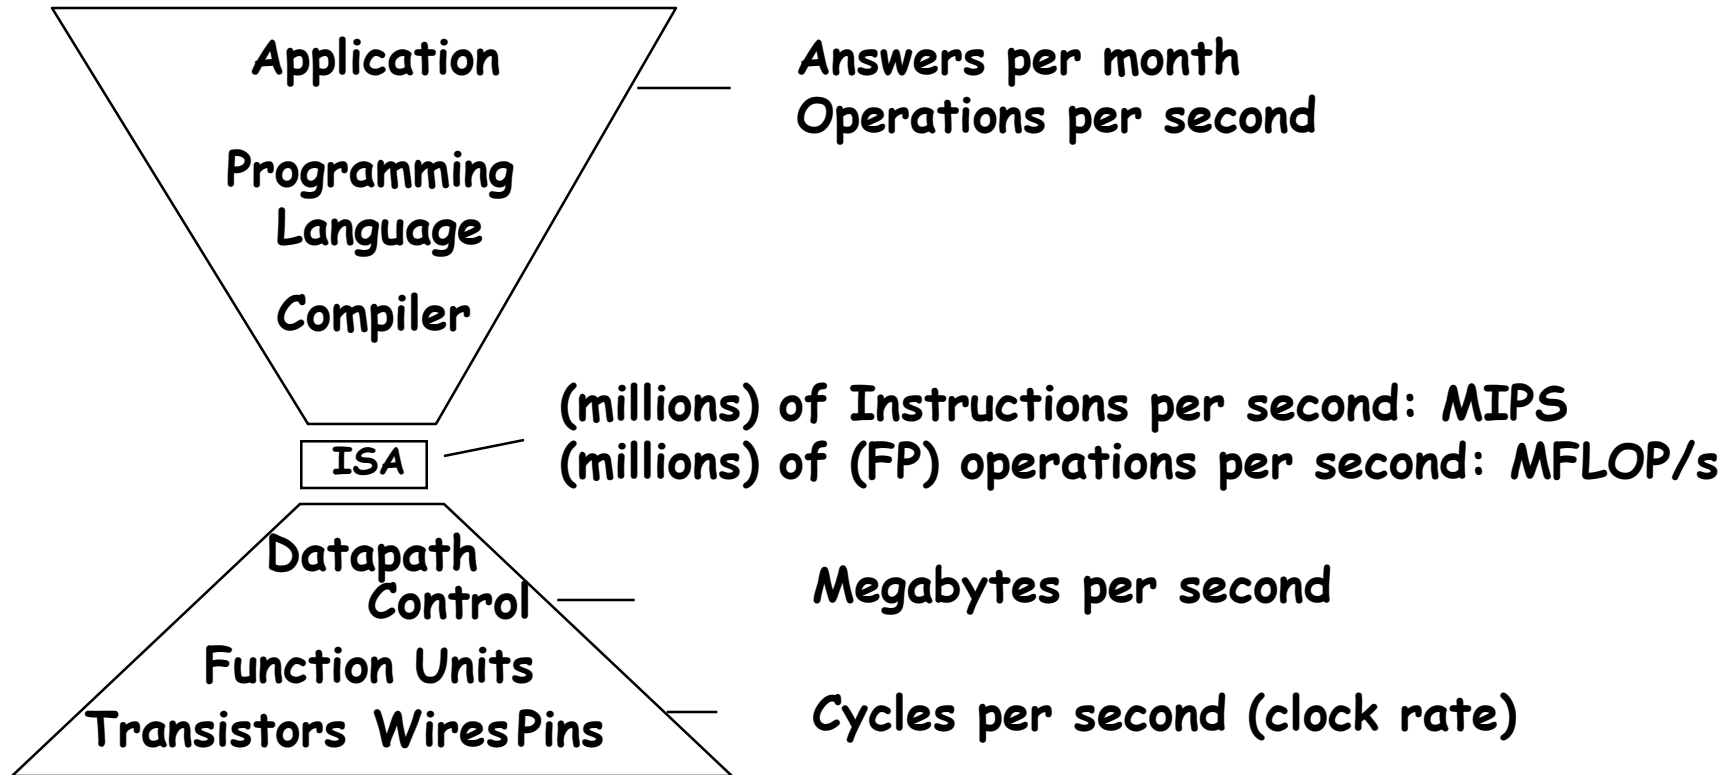


Corollaries

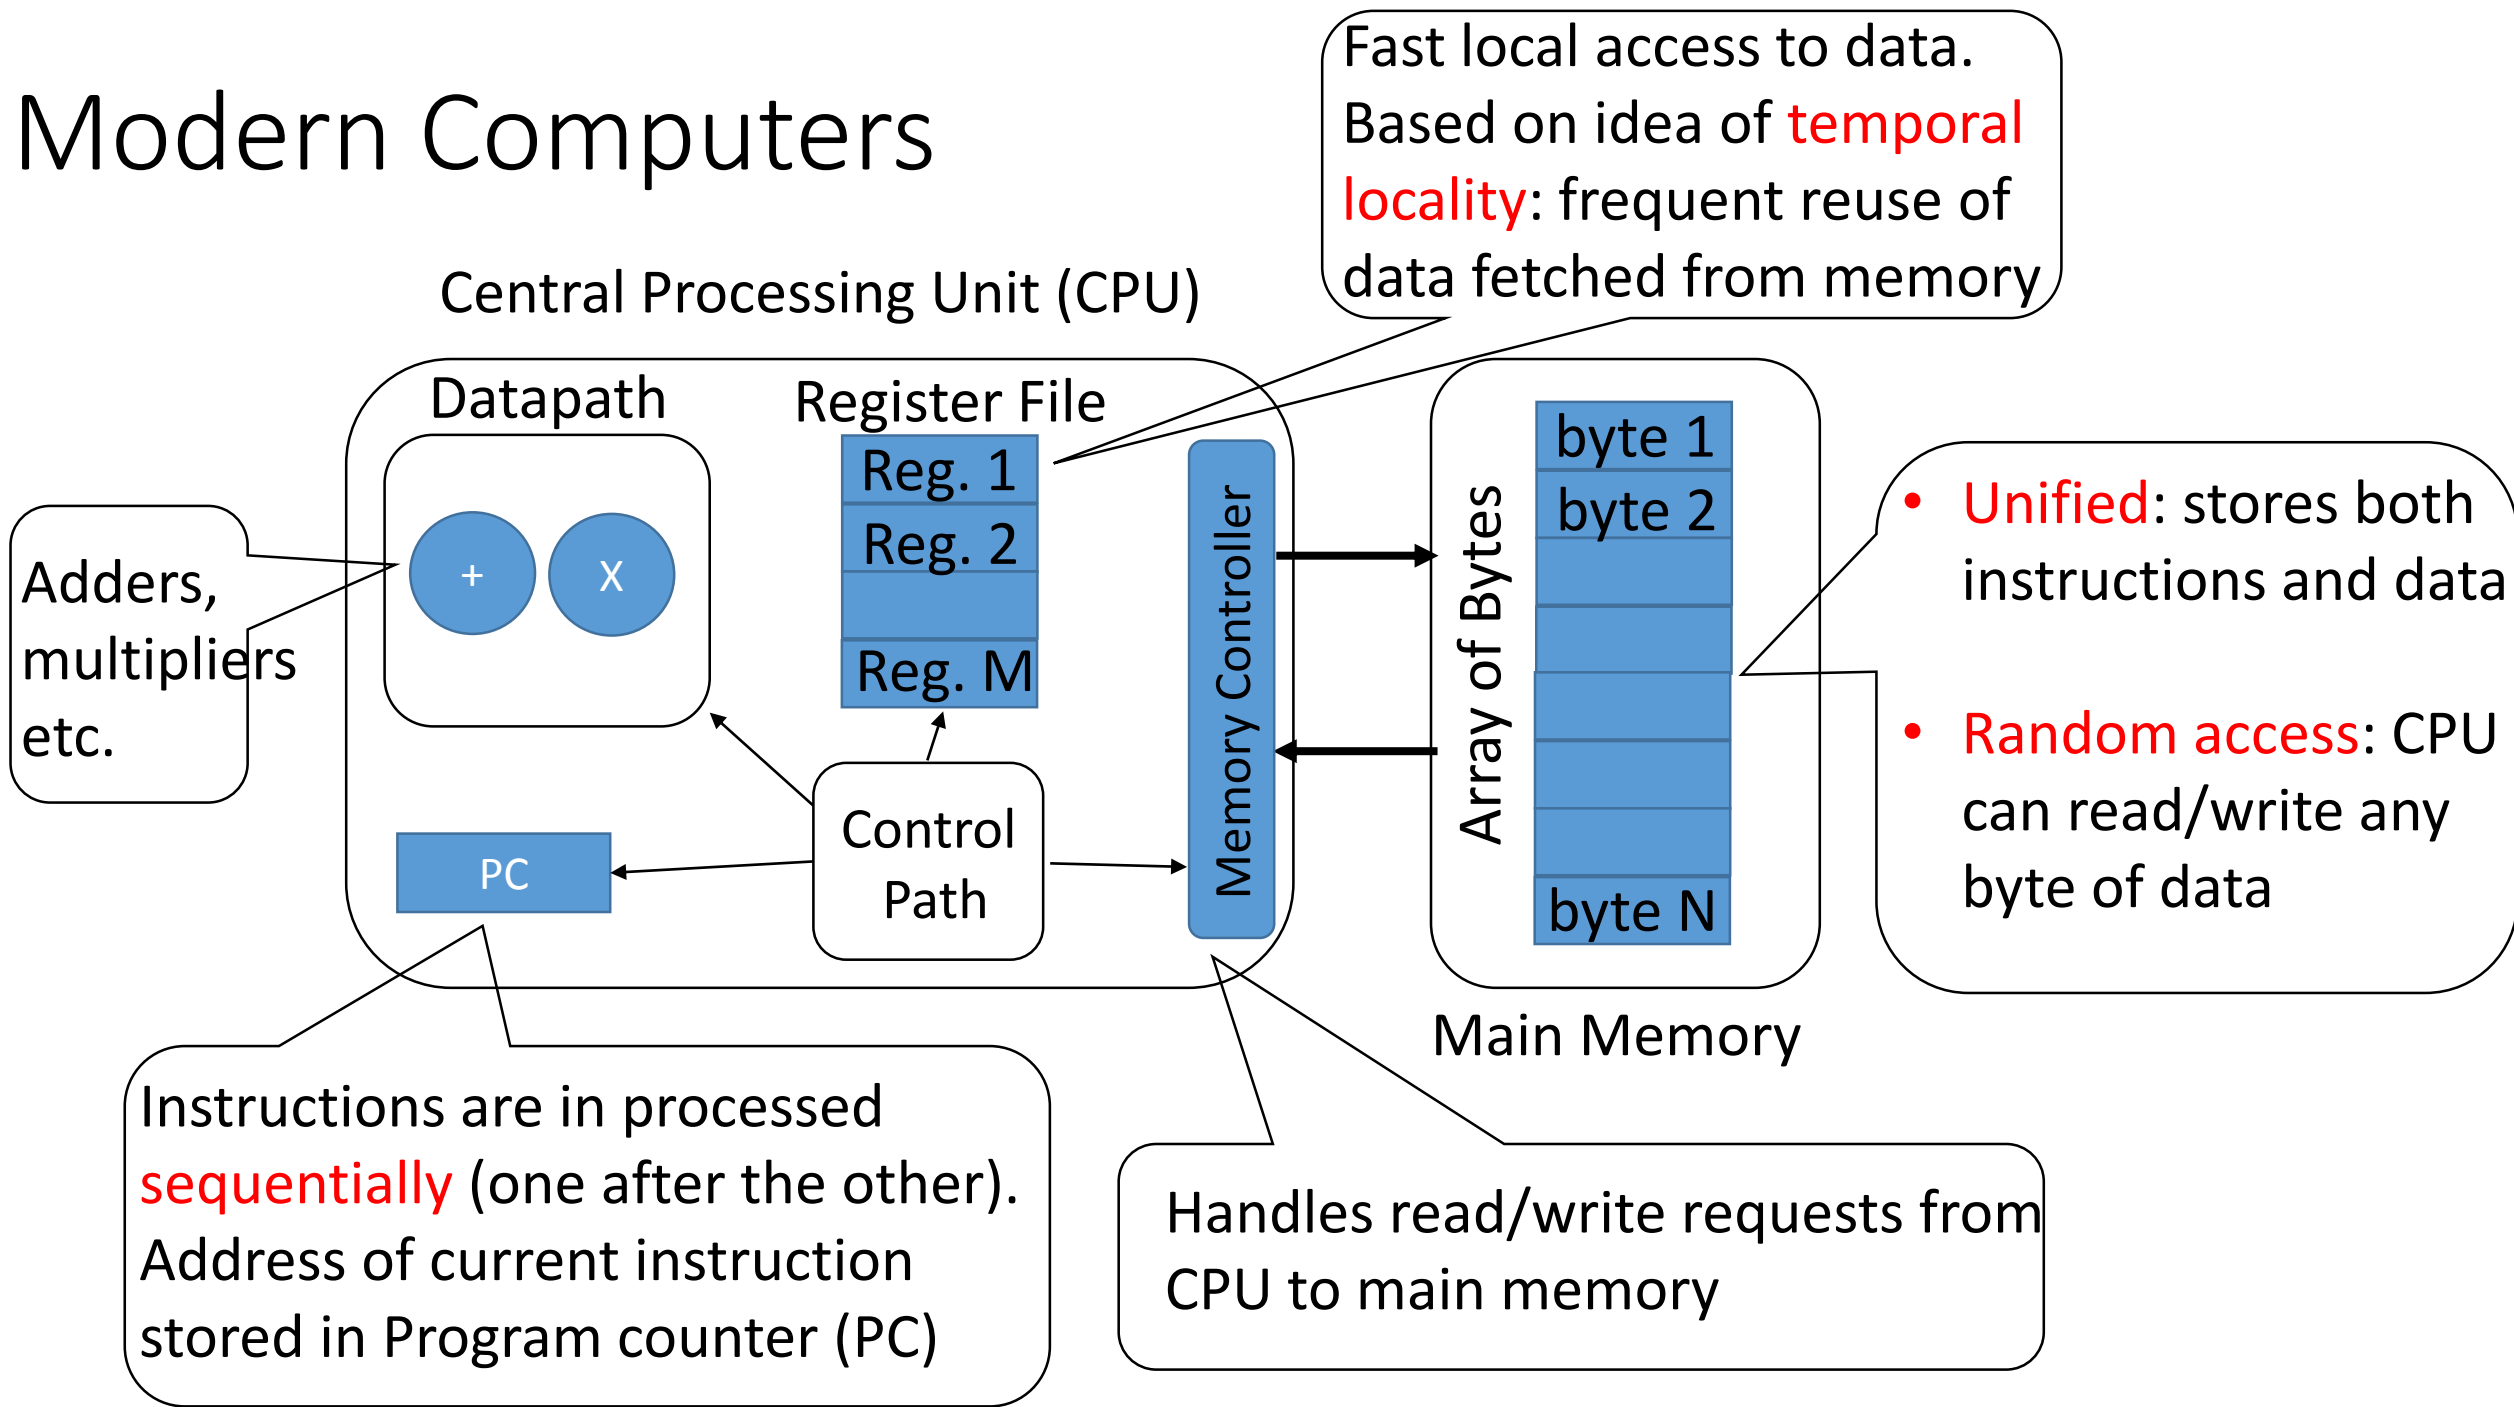
- Many other factors grow exponentially
 - Ex: clock frequency and therefore processor performance



Metrics of Performance



Modern Computers


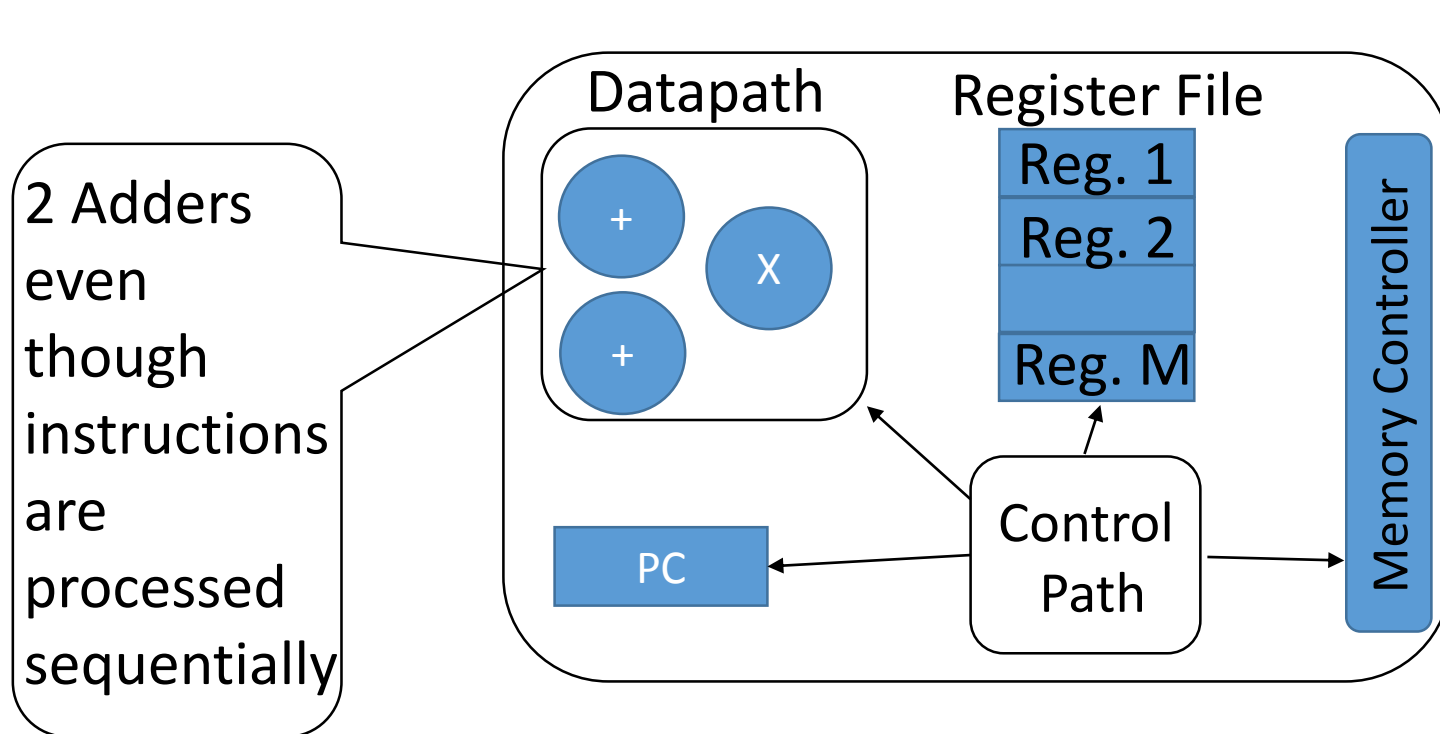


Role of the ISA

- Instruction Set Architecture (ISA) serves as a **contract** between programmer and hardware
 - **Functional description** of storage locations and the operations it can perform
- Storage locations:
 - **Registers**: typically a small number of registers. Ex: MIPS ISA defines 32 32-bit registers.
 - **Memory**: much larger, but off-chip and higher access cost. MIPS ISA has a 32-bit memory address, so in theory 2^{32} Bytes = 4GB of data.
- Operations:
 - Integer Add, Integer Multiply, Branch, FP Add, FP Multiply, FP Divide

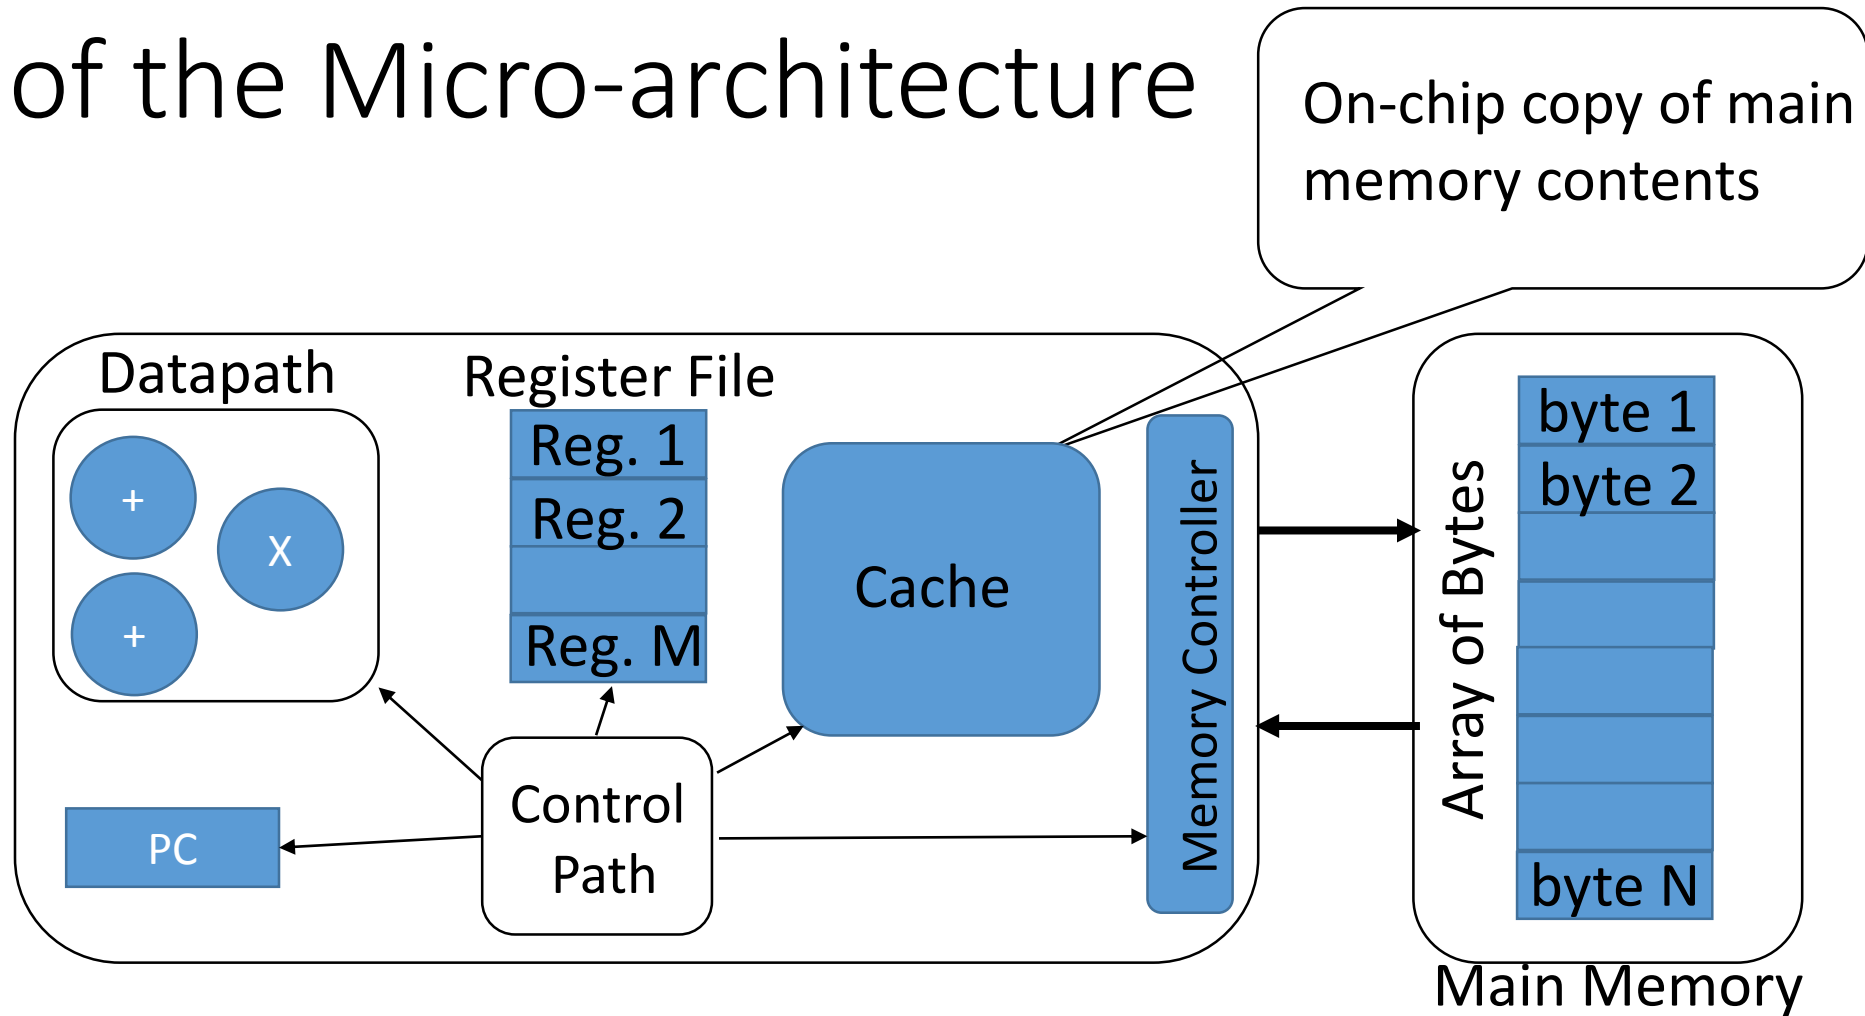
Role of the Micro-architecture

- The micro-architecture is responsible for **faithfully executing** programs written using the ISA
 - Same ISA can have multiple different micro-architectural implementations with different performance and power consumption



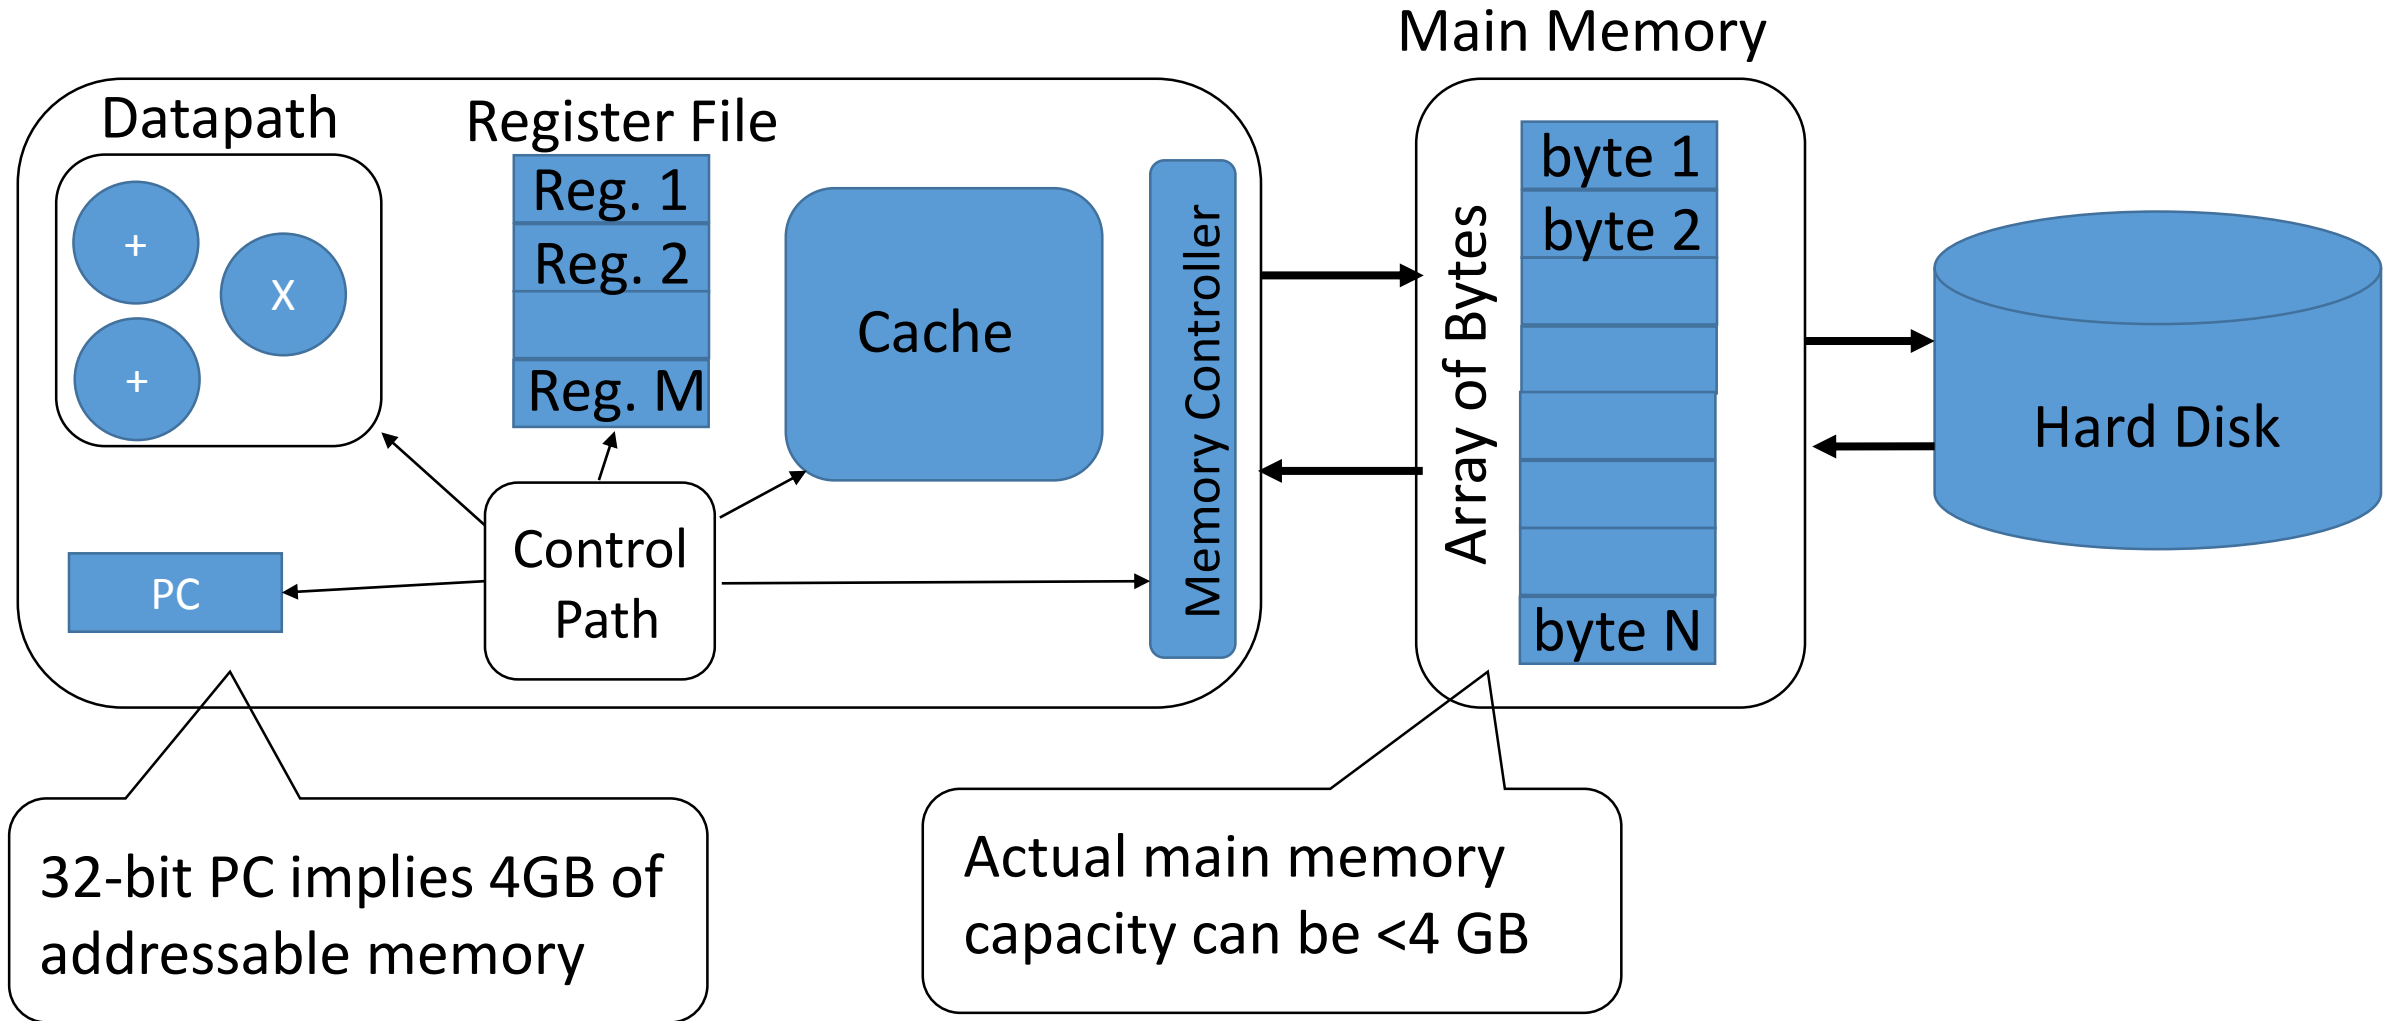
More area/power but
higher performance!
(Identical functionality)

Role of the Micro-architecture



- Cache is managed entirely by hardware
- Programmer and oblivious to its existence





Role of the Micro-architecture




- Hardware/OS managed virtual memory provides the appearance of a 4GB contiguous address space

ISAs are Abundant

RISC ISAs




CISC ISAs



ISA Design: Instruction Length

- An instruction is a basic unit of an ISA
- Fixed versus variable length instructions
 - **Fixed**: Each instruction in the MIPS ISA is 32 bits (4 Bytes) long
 - **Variable**: x86 instructions vary from 1 to 17 Bytes long




Typical of RISC ISAs

Trade-offs?




Typical of CISC ISAs




ISA Design: Operand Addresses

Where do the operands come from?



Where is the result written?



add rs, rt, rd
(\$rd ← \$rs + \$rt)

Value of register rt

sub rs, rt, rd
(\$rd ← \$rs - \$rt)

Instruction specifies 2 source and 1 destination register

"3-Address" ISA (E.g., MIPS ISA)


Example: Swap Using Add/Sub

```
X = X + Y;  
Y = X - Y; // X + Y - Y = X  
X = X - Y; // X + Y - X = Y
```


There are simpler ways to
implement a swap in the MIPS ISA!

- Swap the contents of Reg. 2 (r2) and Reg. 3 (r3) using only *add* and *sub* operations

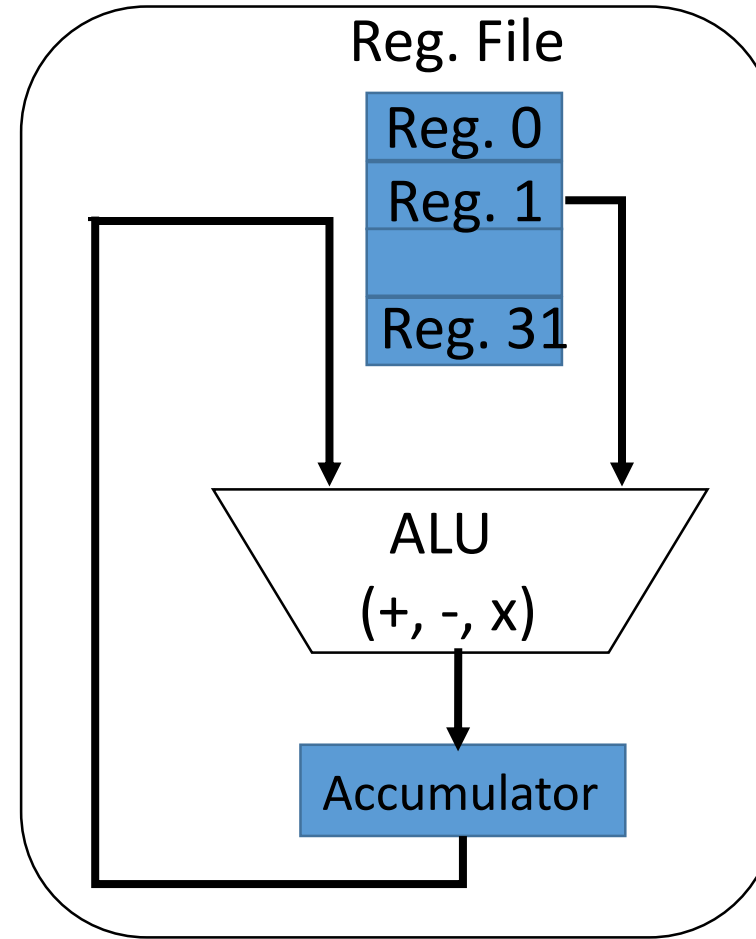
```
add r2, r3, r2  
sub r2, r3, r3  
sub r2, r3, r2
```



1- and 2-Address ISAs




Write result back to a source register
"2-Address" ISA
(Have we seen an example today?)



Special purpose "Accumulator" serves as an operand and as the destination
"1-Address" ISA
(Also called an "Accumulator Architecture")

Swap Routine for 2-Address ISA



add rs, rd
($\$rd \leftarrow \$rd + \$rs$)

sub rs, rd
($\$rd \leftarrow \$rd - \$rs$)

$X = X + Y;$
 $Y = X - Y; \text{ // } X + Y - Y = X$
 $X = X - Y; \text{ // } X + Y - X = Y$


X is stored in r2

Y is stored in r3

Assume r4 can be used as temp

add r3, r2 // $r2 = X + Y$
sub r4, r4 // $r4 = 0$
add r2, r4 // $r4 = X + Y$
sub r3, r4 // $r4 = X$
sub r3, r3 // $r3 = 0$
add r4, r3 // $r3 = X$
sub r3, r2 // $r2 = Y$

Swap Routine for 1-Address ISA



add rs
(\$acc \leftarrow \$acc + \$rs)

sub rs
(\$acc \leftarrow \$acc - \$rs)

mvto rs
(\$rs \leftarrow \$acc)

mvfrm rs
(\$acc \leftarrow \$rs)

$X = X + Y;$
 $Y = X - Y; //$ $X + Y - Y = X$
 $X = X - Y; //$ $X + Y - X = Y$

X is stored in r2

Y is stored in r3

Assume r4 can be used as temp

mvfrm r2 $//acc=X$
add r3 $//acc=X+Y$
mvto r2 $//r2=X+Y$
sub r3 $//acc=X$
mvto r3 $//r3=X$
mvfrm r2 $//acc=X+Y$
sub r3 $//acc=Y$
mvto r2 $//r2=Y$

“Lecture 1: In-class Problem 2” on
newclasses.nyu.edu under “Assignments”

So Which is “Better”?

- Metrics?

- Code size (Bytes)
- Design complexity

```
add r2, r3, r2  
sub r2, r3, r3  
sub r2, r3, r2
```

Assume 5-bit opcode
for all 3 ISAs

How many bits
requires to address
32 registers?

```
add r3, r2  
sub r4, r4  
add r4, r3  
sub r3, r4  
sub r3, r4  
add r3, r2  
sub r2, r3
```


```
mvfrm r2  
add r3  
mvto r2  
sub r3  
mvto r3  
mvfrm r2  
sub r3  
mvto r2
```

- Code size = Num Insts x Bits/Inst


- 3-Address ISA
- 2-Address ISA
- 1-Address ISA

So Which is “Better”?

- Metrics?
 - Code size (Bytes)
 - Design complexity

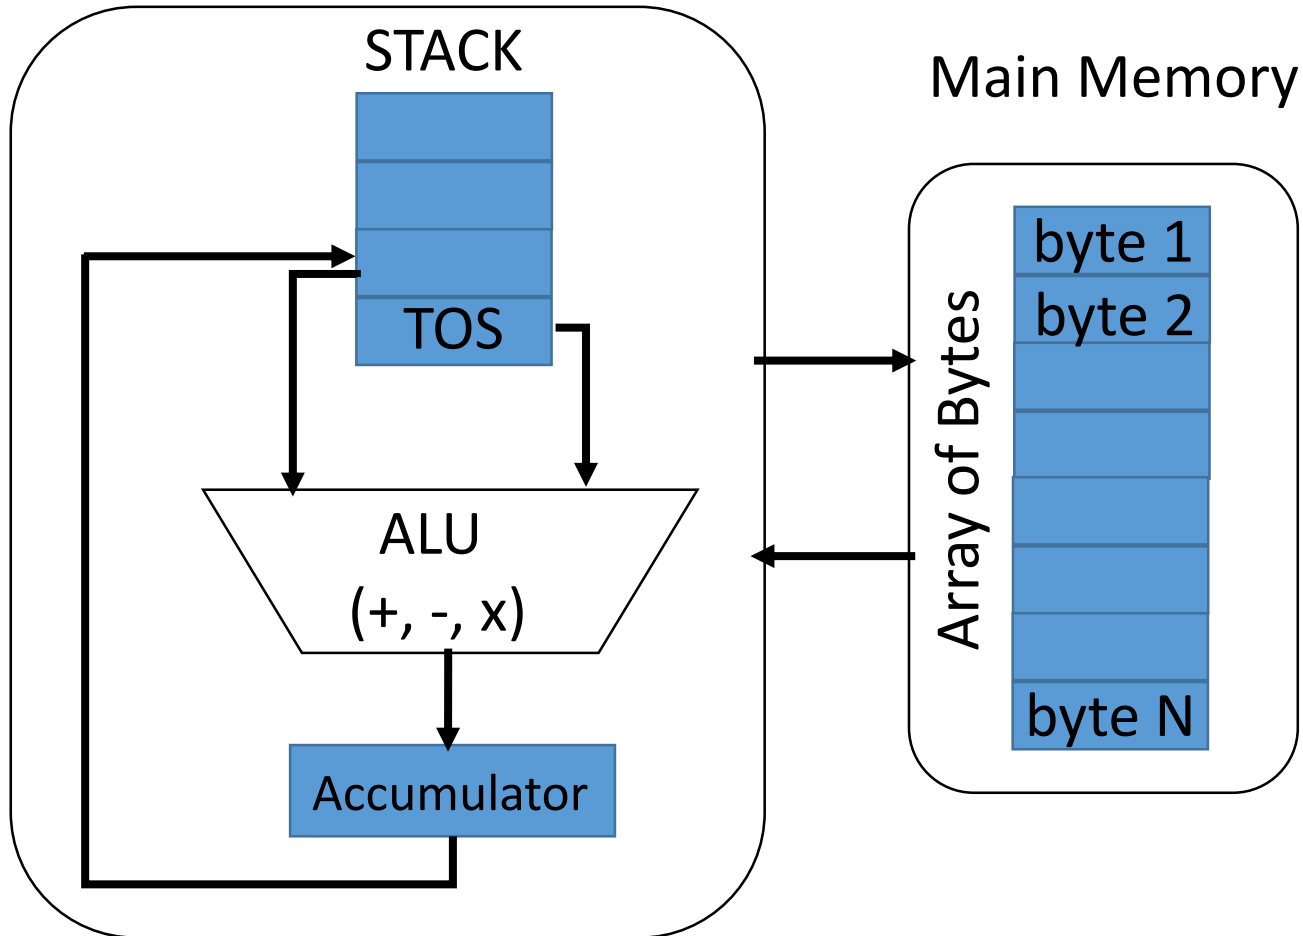


2 Read ports
1 Write port



1 Read ports
+ 1 R/W accum

0-Address ISA (Stack Machines)



pop M

$$\text{Mem}[M] \leftarrow \$TOS$$
$$\text{TOS} \leftarrow \text{TOS}-1$$


push M

$$\text{TOS} \leftarrow \text{TOS}+1$$
$$\$TOS \leftarrow \text{Mem}[M]$$

add

$$\text{\$TOS-1} \leftarrow \text{\$TOS} + \text{\$(TOS-1)}$$
$$\text{TOS} \leftarrow \text{TOS}-1$$

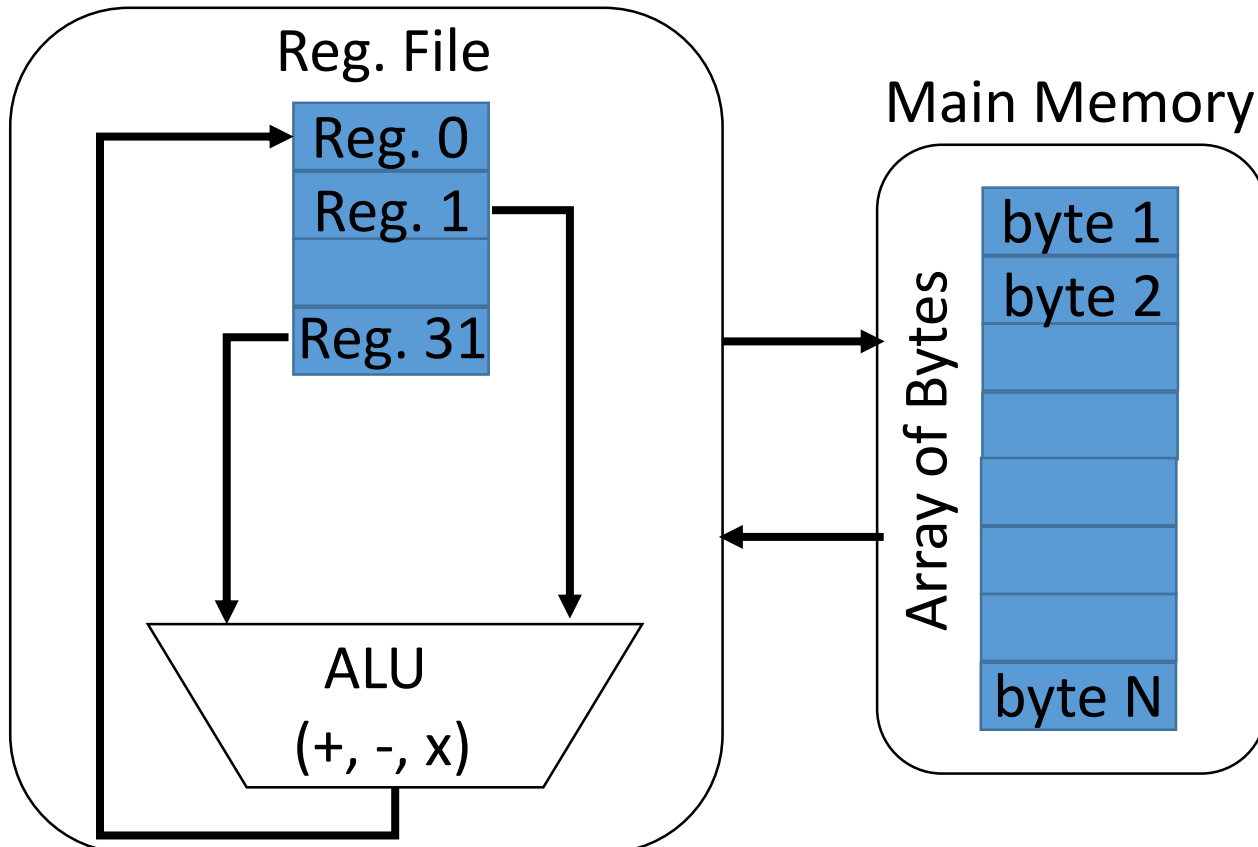
How is Memory Accessed



- Can also be bit-addressable, 32-bit addressable, 64-bit addressable, etc.

Load-Store ISA

- Instructions like *add*, *sub* etc. operate *only* on registers (examples: MIPS, ARM, 3/2/1/0-Address ISAs we have seen so far)
 - Separate Load/Store instructions to fetch data from memory into registers and to write data back from registers to memory



add rs, rt, rd
($\$rd \leftarrow \$rs + \$rt$)

load rt, M
($\$rt \leftarrow \text{Mem}[M]$)

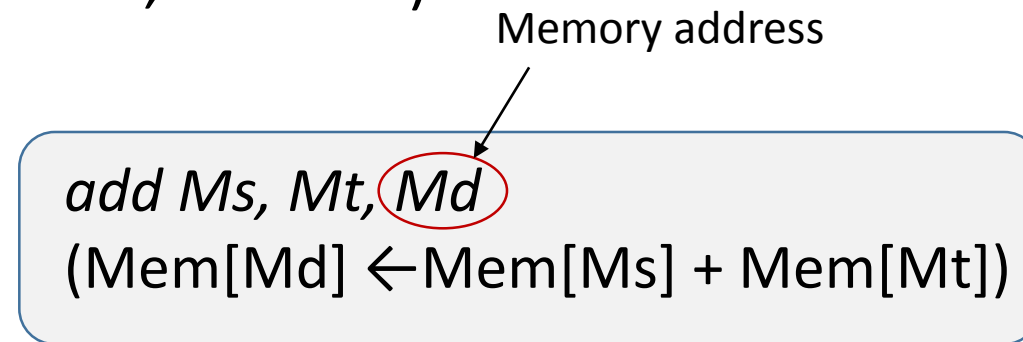
Load data at
memory location M
into register rt

store rt, M
($\text{Mem}[M] \leftarrow \rt)

Store data in register
rt to memory
location M

Memory-Memory ISA

- Allow instructions like add/sub to operate *directly* on data in memory (Examples include: x86, VAX etc.)



- Most generally, each of the two operands and the result can be read-from/written-to either memory or register file
 - Pros?
 - Cons?

Addressing Modes

- How to specify where to read from/write to

Absolute

load rt, M
($\$rt \leftarrow \text{Mem}[M]$)

Directly provide
memory address

Register Indirect

load rs, rt
($\$rt \leftarrow \text{Mem}[\$rs]$)

Memory address
from register rs

Displaced/Based

load rs, rt, offset
($\$rt \leftarrow \text{Mem}[\text{offset} + \$rs]$)

Add an offset to
address stored in
register rs

Indexed

load rs, rt, rd
($\$rd \leftarrow \text{Mem}[\$rs + \$rt]$)

Base memory
address and offset
from registers

Memory Indirect

load rt, rs
($\$rt \leftarrow \text{Mem}[\text{Mem}[\$rs]]$)

Use value at
 $\text{Mem}[\$rs]$ as
address

Auto-increment

load rt, rs
($\$rt \leftarrow \text{Mem}[\$rs]$
 $\$rs \leftarrow \$rs + 1$)

Same as Register
indirect + value in
rs is automatically
incremented

- Orthogonal ISA:** each op supports every addressing mode (ex: VAX)

What Makes a Good ISA

- Ease of hardware implementations
 - Decoder for fixed vs. variable instruction lengths
 - Simple vs. complex addressing modes
- Ease of software implementation
 - Can the programmer/compiler use the ISA easily?
 - How many assembly instructions to represent a single line of code in a high-level language? “semantic gap”
- Backwards compatibility
 - ISA will have to be supported into the future; new instructions can be added but existing instructions cannot be removed

CISC Vs. RISC

- CISC: “Complex” Instruction Set Computer
- RISC: “Reduced” Instruction Set Computer
 - Each CISC instruction = multiple RISC instructions
- CISC vs. RISC debate dominated the 80s/early 90s but is largely resolved now
 - CISC ISAs like x86 are “broken down” internally into simpler RISC-like instructions (microcode)
 - x86 is “externally CISC but internally RISC”
 - Other popular ISAs like ARM and MIPS are RISC

ISA Comparison

	VAX	x86	MIPS
ISA Type	CISC	CISC (but internally RISC)	RISC
Instruction Length	Variable	Variable	Fixed
Addressing Modes	3-Address; memory-memory; orthogonal addressing modes	2-Address; register-memory	3-Address; load/store ISA