

Computer Architecture I


CS-GY-6133

Topic: Introduction and Instruction Set Architectures (ISA)


Instructor: Siddharth Garg

ISA Design: Operand Addresses

Where do the operands come from?



Where is the result written?



add rs, rt, rd
(\$rd \leftarrow \$rs + \$rt)


Value of register rt

sub rs, rt, rd
(\$rd \leftarrow \$rs - \$rt)

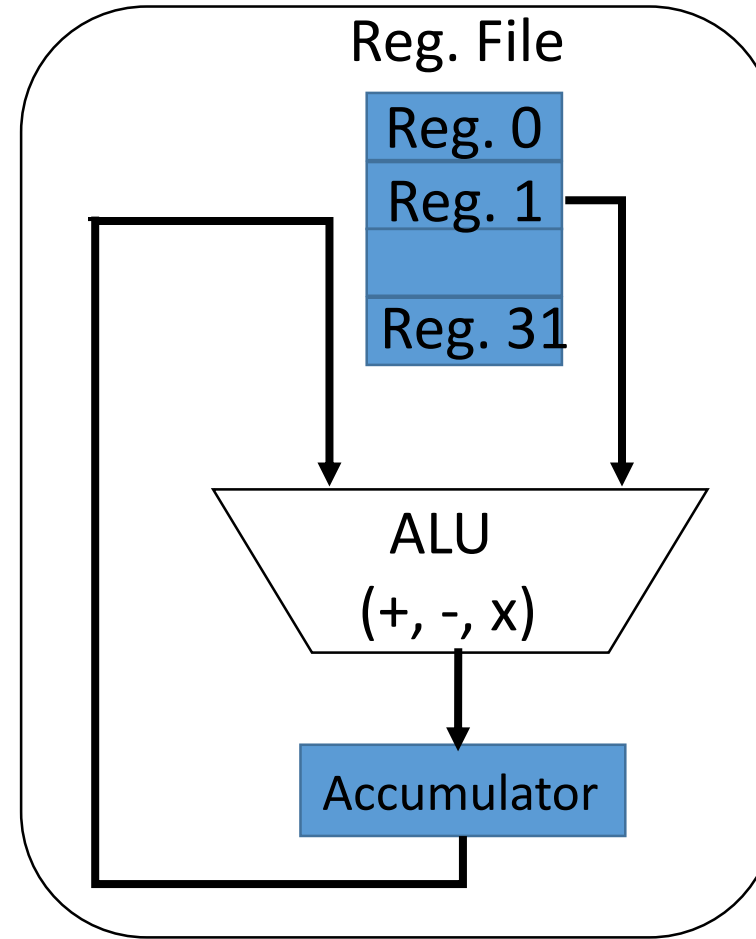
Instruction specifies 2 source and 1 destination register

"3-Address" ISA (E.g., MIPS ISA)

1- and 2-Address ISAs



Write result back to a source register
"2-Address" ISA
(Have we seen an example today?)



Special purpose "Accumulator" serves as an operand and as the destination
"1-Address" ISA
(Also called an "Accumulator Architecture")

So Which is “Better”?

- Metrics?

- Code size (Bytes)
- Design complexity

```
add r2, r3, r2  
sub r2, r3, r3  
sub r2, r3, r2
```

Assume 5-bit opcode
for all 3 ISAs

How many bits
requires to address
32 registers?

```
add r3, r2  
sub r4, r4  
add r4, r3  
sub r3, r4  
sub r3, r4  
add r3, r2  
sub r2, r3
```


```
mvfrm r2  
add r3  
mvto r2  
sub r3  
mvto r3  
mvfrm r2  
sub r3  
mvto r2
```

- Code size = Num Insts x Bits/Inst


- 3-Address ISA
- 2-Address ISA
- 1-Address ISA

So Which is “Better”?

- Metrics?
 - Code size (Bytes)
 - Design complexity

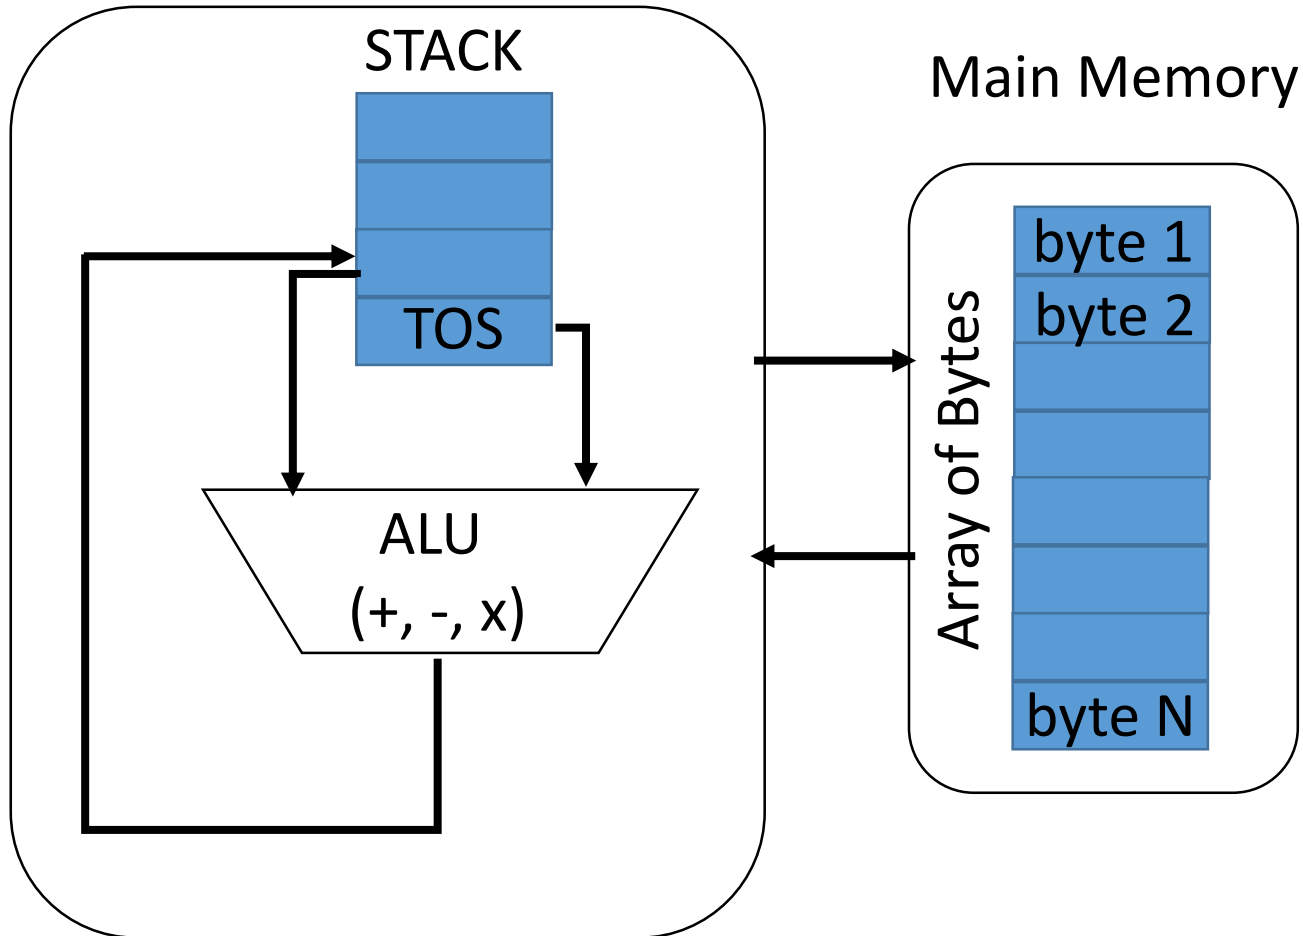


2 Read ports
1 Write port



1 Read ports
+ 1 R/W accum

0-Address ISA (Stack Machines)



pop M

$\text{Mem}[M] \leftarrow \TOS

$\text{TOS} \leftarrow \text{TOS}-1$

push M

$\text{TOS} \leftarrow \text{TOS}+1$

$\$\text{TOS} \leftarrow \text{Mem}[M]$

add

$\$\text{TOS}-1 \leftarrow \$\text{TOS} + \$(\text{TOS}-1)$

//Pull out two operands from stack,
//replace with result

$\text{TOS} \leftarrow \text{TOS}-1$

//Update stack pointer

In-Class Problem 1

Variables x, y, w, z are stored at addresses $\&x, \&y, \&w$ and $\&z$ in main memory.

Write stack machine code to compute:

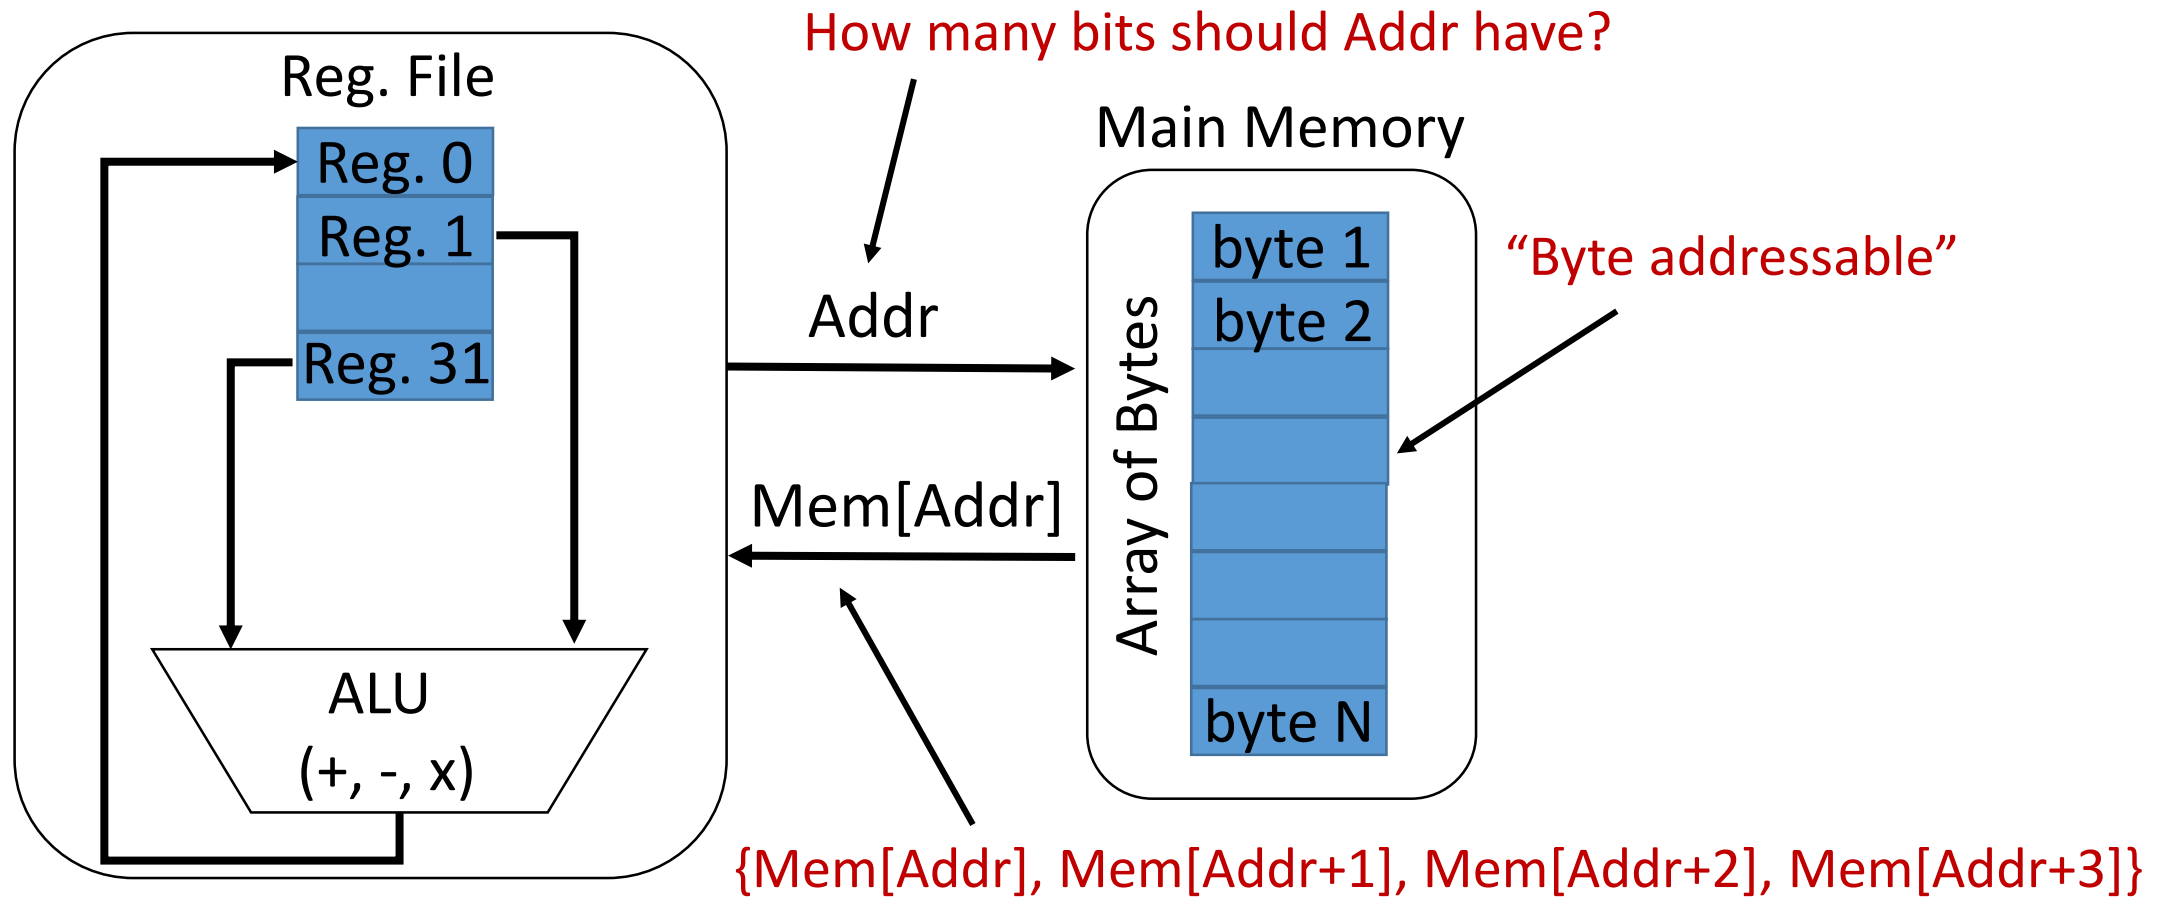
$$f = ((x * y) + w) * z$$

and store the result in memory location $\&f$.

Assume that the *add* and *mult* operations are available.

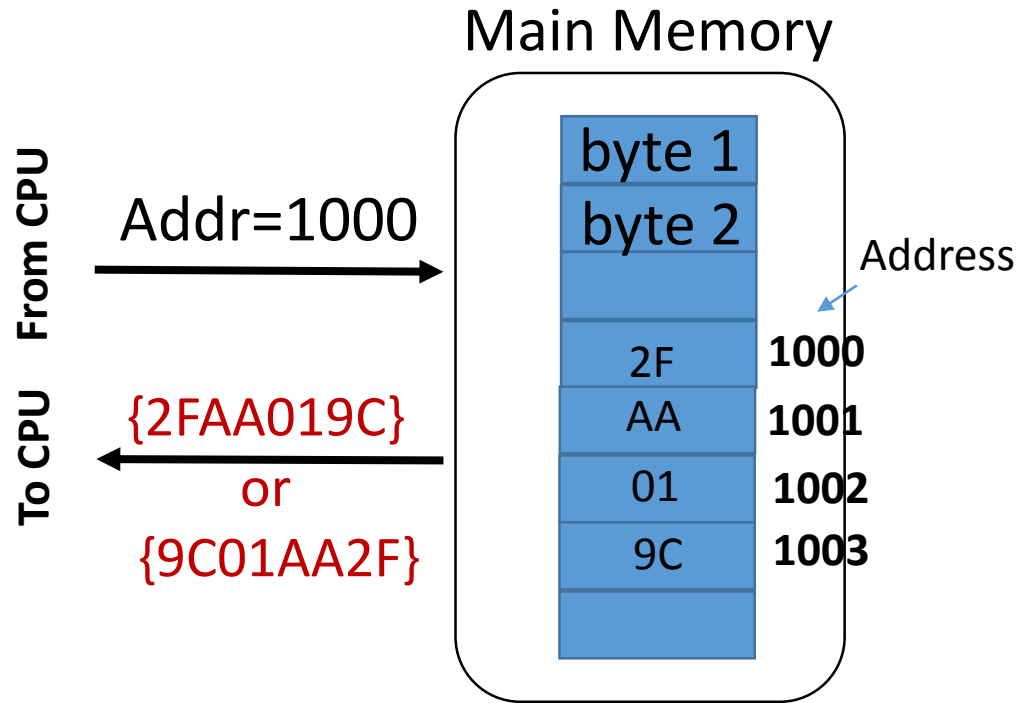
```
push &x  
push &y  
mult // $TOS = x*y  
push &w  
add // $TOS = (x*y) + w  
push &z  
mult // $TOS = ((x*y) + w)*z  
pop &f
```

How is Memory Accessed



- Memory can also be bit-addressable, 32-bit addressable, 64-bit addressable, etc.

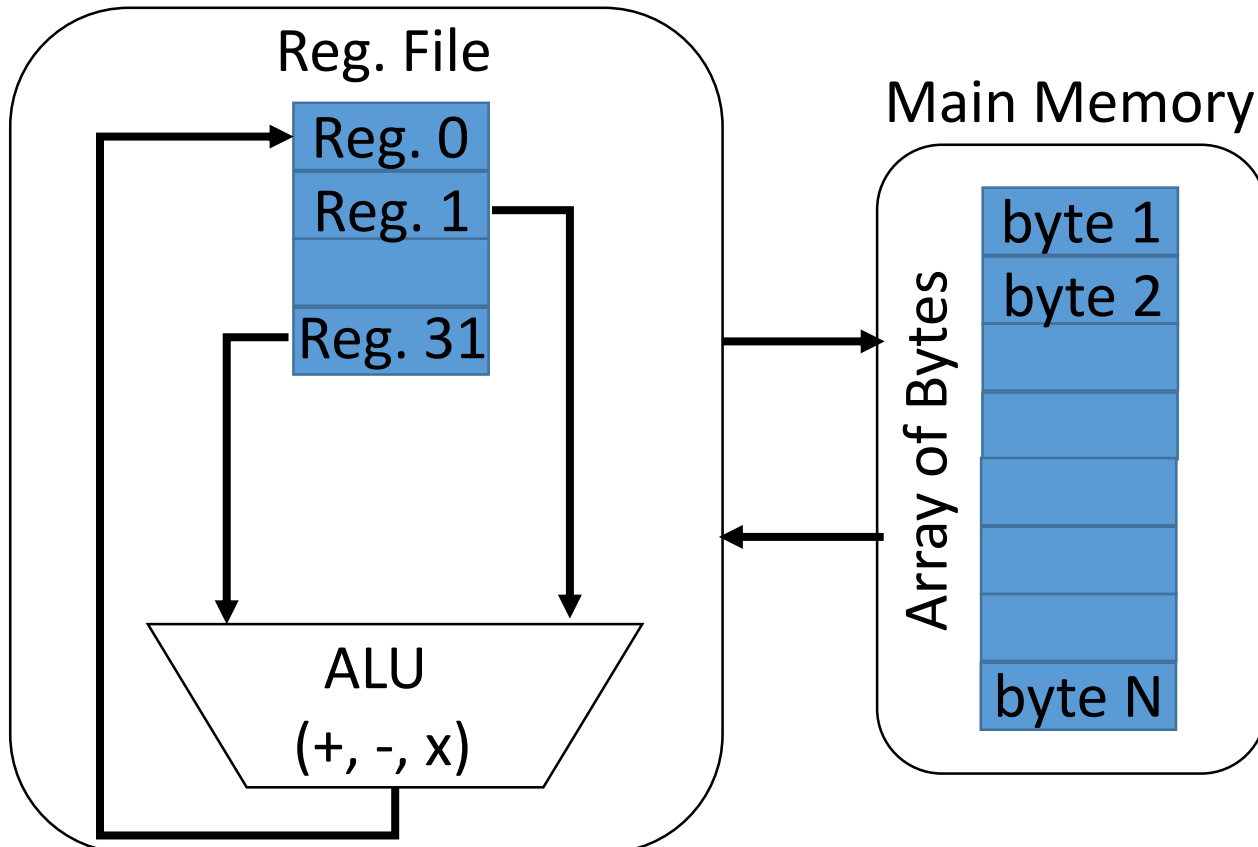
Little/Big Endian



- Assuming an instruction or data word is 32 bits of data from locations [Addr,Addr+3] are fetched from Main Memory.
- Big Endian:** Data = {Mem[Addr], Mem[Addr+1], Mem[Addr+2], Mem[Addr+3]}
 - Examples: IBM360
- Little Endian:** Data = {Mem[Addr+3], Mem[Addr+2], Mem[Addr+1], Mem[Addr]}
 - Examples: x86, x86-64, AMD64
- Bi-Endian: both supported! (Ex: MIPS)

Load-Store ISA

- Instructions like *add*, *sub* etc. operate *only* on registers (examples: MIPS, ARM, 3/2/1/0-Address ISAs we have seen so far)
 - Separate Load/Store instructions to fetch data from memory into registers and to write data back from registers to memory



add rs, rt, rd
($\$rd \leftarrow \$rs + \$rt$)

load rt, M
($\$rt \leftarrow \text{Mem}[M]$)

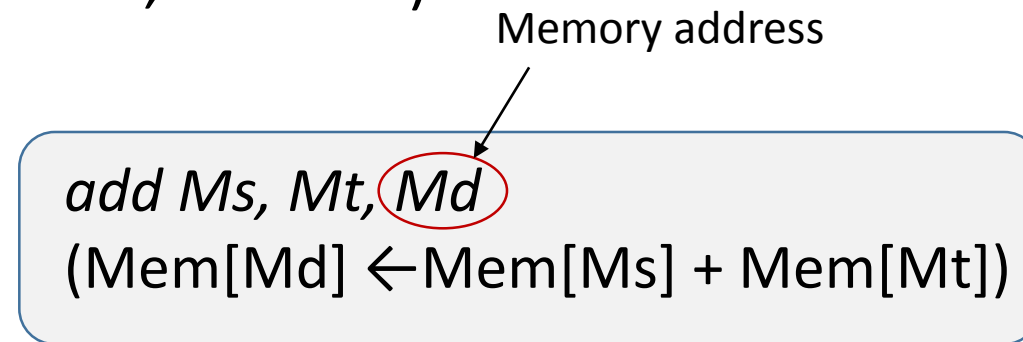
Load data at
memory location M
into register rt

store rt, M
($\text{Mem}[M] \leftarrow \rt)

Store data in register
rt to memory
location M

Memory-Memory ISA

- Allow instructions like add/sub to operate *directly* on data in memory (Examples include: x86, VAX etc.)



- Most generally, each of the two operands and the result can be read-from/written-to either memory or register file
 - Pros?
 - Cons?

Addressing Modes

- How to specify where to read from/write to

Absolute

load rt, M
 $(\$rt \leftarrow \text{Mem}[M])$

Directly provide
memory address

Register Indirect

load rs, rt
 $(\$rt \leftarrow \text{Mem}[\$rs])$

Memory address
from register rs

Displaced/Based

load rs, rt, offset
 $(\$rt \leftarrow \text{Mem}[\text{offset} + \$rs])$

Add an offset to
address stored in
register rs

Indexed

load rs, rt, rd
 $(\$rd \leftarrow \text{Mem}[\$rs + \$rt])$

Base memory
address and offset
from registers

Memory Indirect

load rt, rs
 $(\$rt \leftarrow \text{Mem}[\text{Mem}[\$rs]])$

Use value at
 $\text{Mem}[\$rs]$ as
address

Auto-increment

load rt, rs
 $(\$rt \leftarrow \text{Mem}[\$rs])$
 $\$rs \leftarrow \$rs + 1$

Same as Register
indirect + value in
rs is automatically
incremented

- Orthogonal ISA:** each op supports every addressing mode (ex: VAX)

Addressing Modes

Absolute

load rt, M
($\$rt \leftarrow \text{Mem}[M]$)

32-bit instruction
and can't possible
store a 32-bit
address!

Register Indirect

load rs, rt
($\$rt \leftarrow \text{Mem}[\$rs]$)

1. Pointers
2. Emulate
absolute
addressing
(more later)

Displaced/Based

load rs, rt, offset
($\$rt \leftarrow \text{Mem}[\text{offset} + \$rs]$)

Data is frequently
stored *relative* to
a base location

Indexed

load rs, rt, rd
($\$rd \leftarrow \text{Mem}[\$rs + \$rt]$)

Memory Indirect

load rt, rs
($\$rt \leftarrow \text{Mem}[\text{Mem}[\$rs]]$)

Pointer to a
pointer!

Auto-increment

load rt, rs
($\$rt \leftarrow \text{Mem}[\$rs]$
 $\$rs \leftarrow \$rs + 1$)

Array accesses
(rs stores pointer
to an array)

What Makes a Good ISA

- Ease of hardware implementations
 - Decoder for fixed vs. variable instruction lengths
 - Simple vs. complex addressing modes
- Ease of software implementation
 - Can the programmer/compiler use the ISA easily?
 - How many assembly instructions to represent a single line of code in a high-level language? “semantic gap”
- Backwards compatibility
 - ISA will have to be supported into the future; new instructions can be added but existing instructions cannot be removed

CISC Vs. RISC

- CISC: “Complex” Instruction Set Computer
- RISC: “Reduced” Instruction Set Computer
 - Each CISC instruction = multiple RISC instructions
- CISC vs. RISC debate dominated the 80s/early 90s but is largely resolved now
 - CISC ISAs like x86 are “broken down” internally into simpler RISC-like instructions (microcode)
 - x86 is “externally CISC but internally RISC”
 - Other popular ISAs like ARM and MIPS are RISC

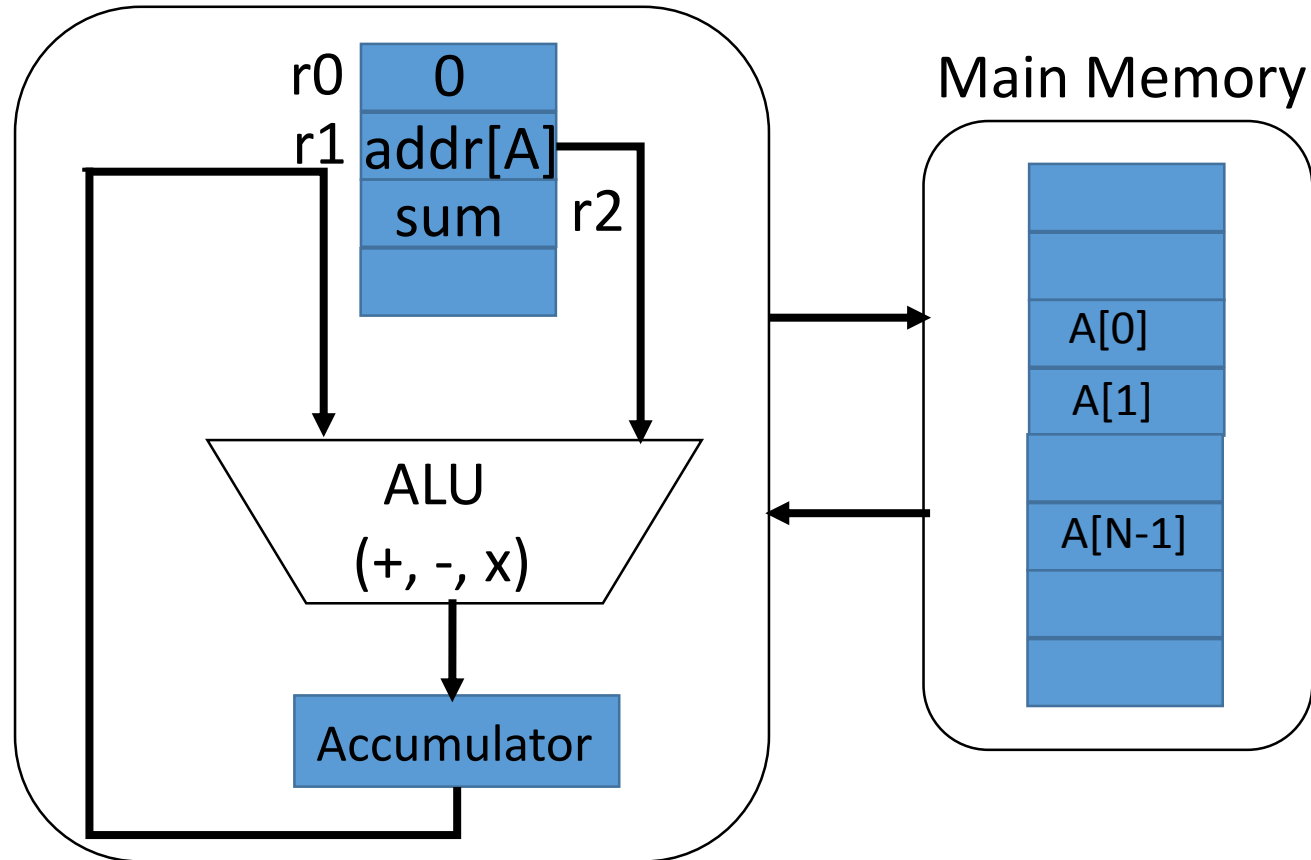
ISA Comparison

	VAX	x86	MIPS
ISA Type	CISC	CISC (but internally RISC)	RISC
Instruction Length	Variable	Variable	Fixed (32 bits)
Addressing Modes	<ul style="list-style-type: none">• 3-Address• Memory-memory• Orthogonal addressing modes	<ul style="list-style-type: none">• 2-address• Register-memory	<ul style="list-style-type: none">• 3-Address• Load/store ISA

In-Class Problem 2

We will explore the advantages of

- (i) The auto-increment addressing mode
- (ii) Orthogonal ISA



mvto rs
($\$rs \leftarrow \acc)

mvfrm rs
($\$acc \leftarrow \rs)

add rs
($\$acc \leftarrow \$acc + \$rs$)

addinc rs, rt
($\$acc \leftarrow \$acc + \text{Mem}[\$rs + \$rt]$
 $\$rs \leftarrow \$rs + 1$)

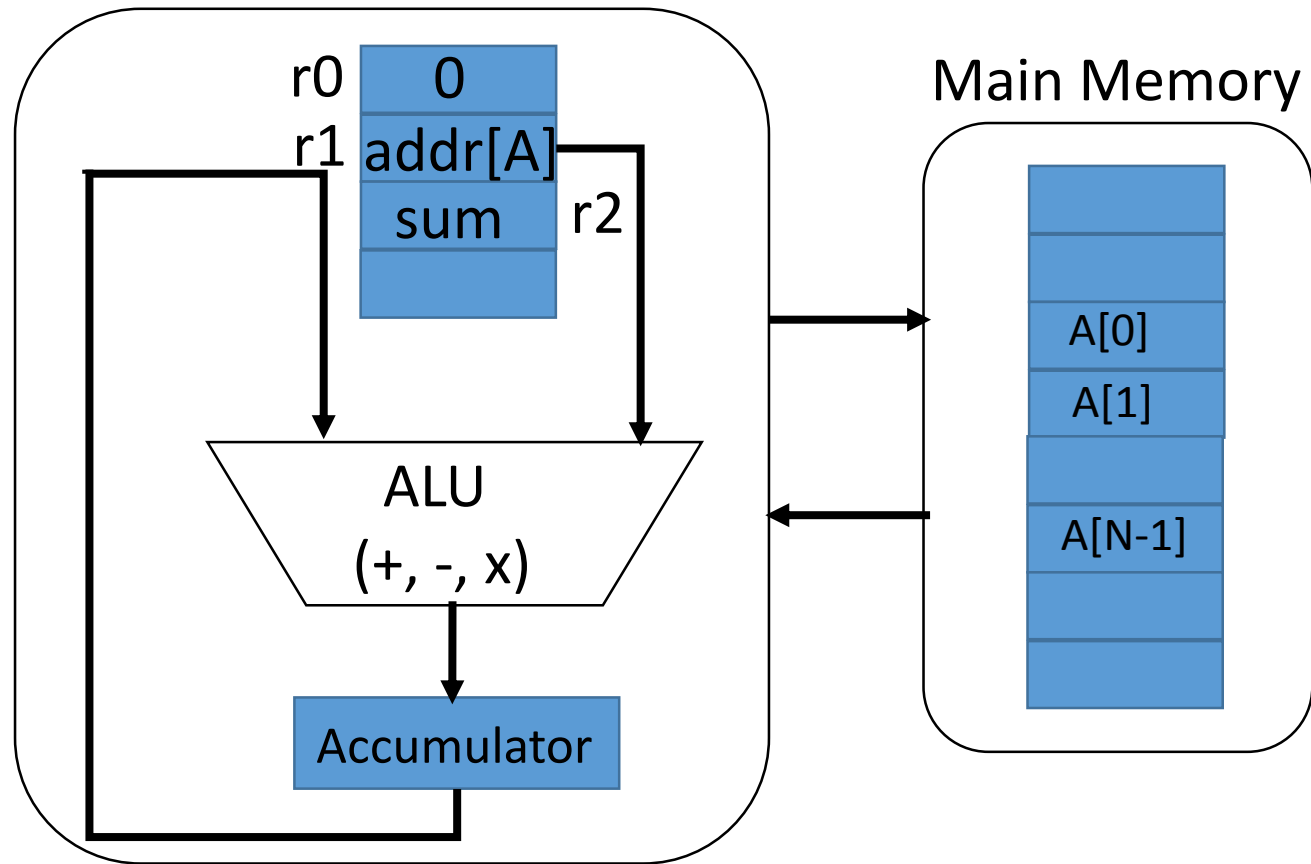
ble rs, imm, addr
(if ($\$rs \leq \text{imm}$)
goto addr)

Write code to sum the array and store the sum in register r2. You can assume the acc is initialized to zero.

In-Class Problem 2

We will explore the advantages of

- (i) The auto-increment addressing mode
- (ii) Orthogonal ISA

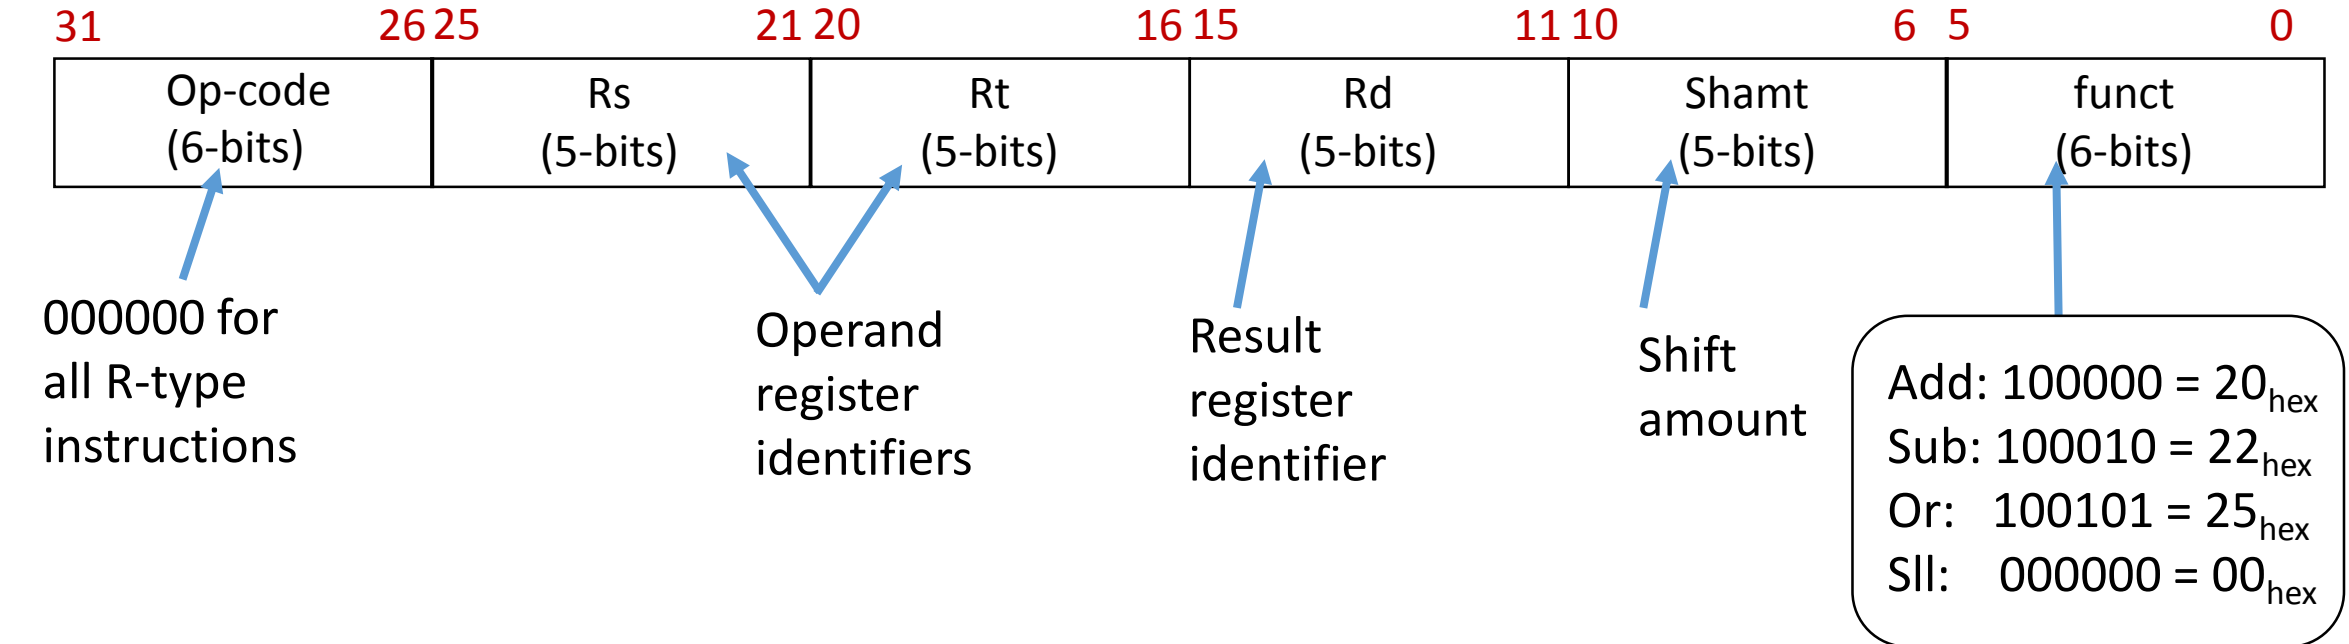


```
mvfrm r0  
L: addinc r0, r1  
ble r0, N-1, L  
mvto r2
```

MIPS ISA Overview

- 32-bit ISA (what does this mean?)
 - **Instruction length**: 32 bits
 - **Data word length**: 32 bits
- 32 general purpose registers (\$0, \$1, \$2, ... \$31); 32-bits each
 - Register \$0 is hard-wired to 0 (constant)
 - All other registers can be used as read/write
 - But, certain registers dedicated for specific purposes by compiler
 - Global pointer (\$28), Stack pointer (\$29), Frame pointer (\$30), Return Address (\$31)
 - We won't worry about these conventions in this class!
- Two dedicated 32-bit registers *hi* and *lo* for multiplication ops
- Separate regs/unit for floating point ops (discussed later)

MIPS Instructions: R-Type



add rs, rt, rd // $R[rd] \leftarrow R[rs] + R[rt]$; signed addition; trap on overflow
sub rs, rt, rd // $R[rd] \leftarrow R[rs] - R[rt]$; signed subtraction; trap on overflow
or rs, rt, rd // $R[rd] \leftarrow R[rs] \mid R[rt]$; bit-wise Boolean OR operation
sll rt, rd, shamt // $R[rd] \leftarrow R[rt] \ll \text{shamt}$; logical shift left
.....(many others)

MIPS R-type Instructions: Add vs Addu

`add rs, rt, rd` // $R[rd] \leftarrow R[rs] + R[rt]$; signed addition; trap on overflow

- rs, rt and rd are assumed to contain 32-bit signed, 2's complement numbers
 - MSB (32nd bit) indicates if the number is positive or negative
 - Review 2's complement representation (https://en.wikipedia.org/wiki/Two%27s_complement)
- Example: assume only 4-bit numbers in range [-8,7]

rs = 0101 (5)

rt =+ 0110 (6)

rd= 1011 (-4) → **Overflow!**


`addu rs, rt, rd` // $R[rd] \leftarrow R[rs] + R[rt]$;

NO trap on overflow

MIPS R-type Instructions: Mult

mult rs, rt // $\{hi, lo\} \leftarrow R[rs] \times R[rd]$; signed multiplication

- *hi* and *lo* are two *special purpose 32-bit registers* that store the result of a multiplication operation
 - Why? Multiplication result of two 32-bit values needs at least 64 bits




What does this remind you of?

mvhi rd // $R[rd] \leftarrow R[hi]$

mvlo rd // $R[rd] \leftarrow R[lo]$

MIPS Instructions: I-Type



addi: 8_{hex}
 beq: 4_{hex}
 ori: 25_{hex}
 lw: 23_{hex}

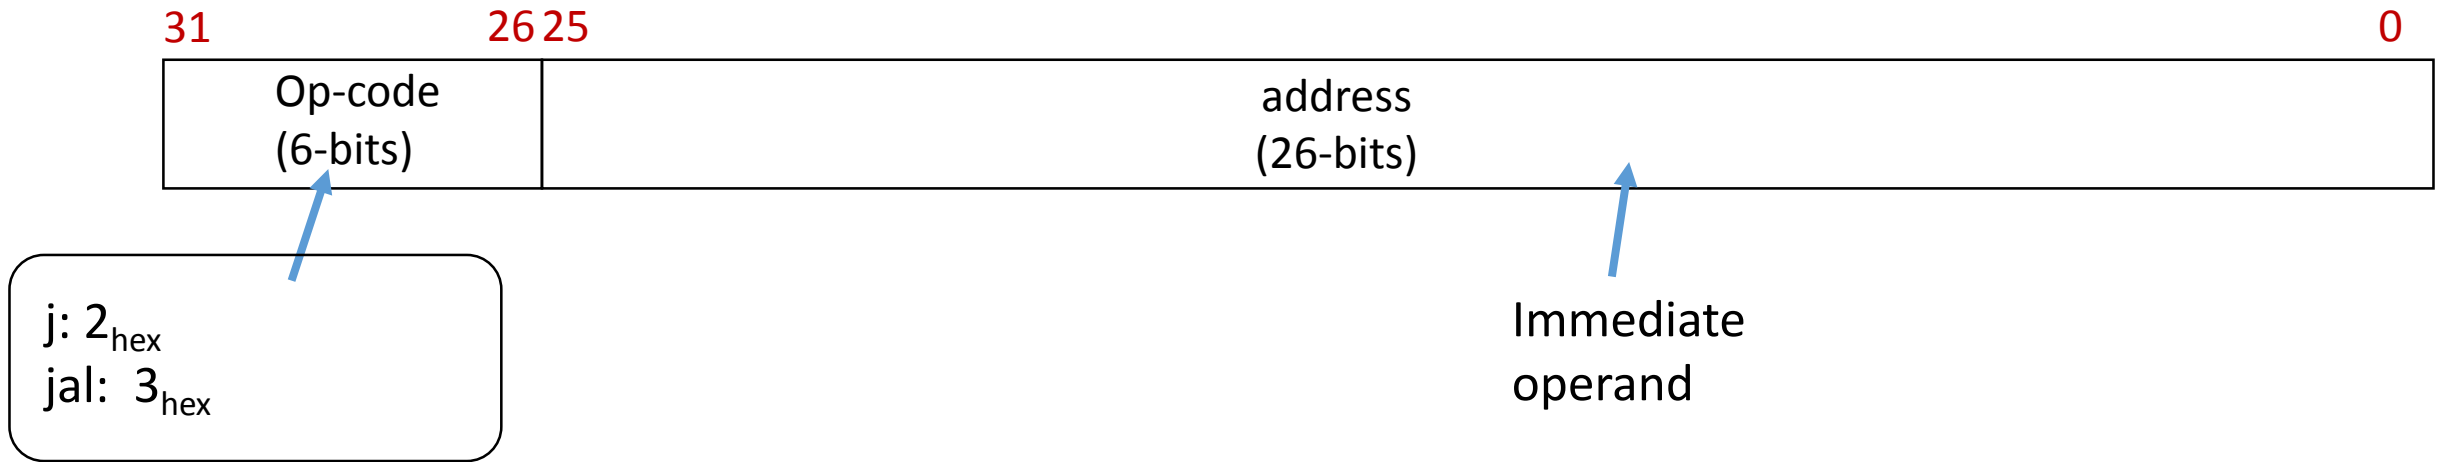
Operand
register

Result
register ("PC relative")

Immediate
operand

addi rs, rt, imm // $R[rt] \leftarrow R[rs] + \{\text{SignExtend}, \text{imm}\}$; MSB of imm is extended to 32 bits
 beq rs, rt, imm // if $\{R[rs] == R[rt]\}$ branch to $PC + 4 + \text{imm}$; ("PC relative")
 ori rs, rt, imm // $R[rt] \leftarrow R[rs] | \{\text{ZeroExtend}, \text{imm}\}$; bit-wise Boolean OR operation
 lw rs, rt, imm // $R[rt] \leftarrow \text{Mem}[\{\text{SignExtendimm}\} + R[rs]]$ ("Displaced/based")
(many others)

MIPS Instructions: J-Type



j address // PC \leftarrow address

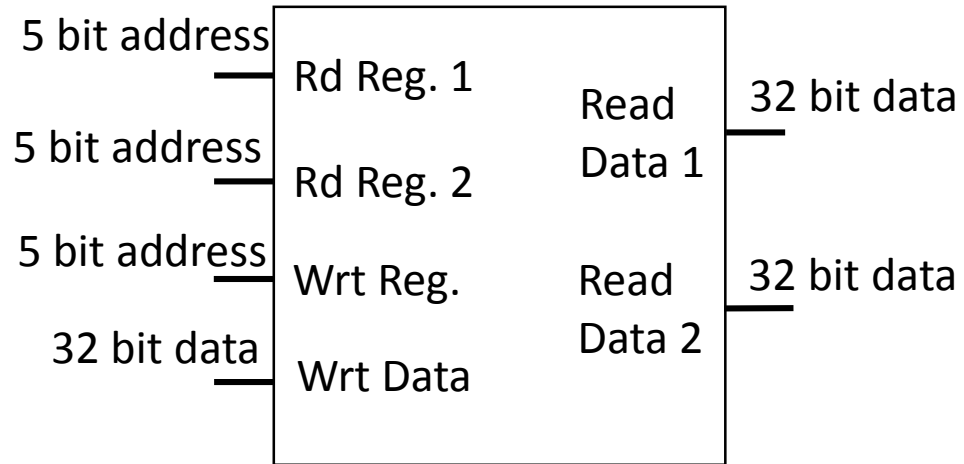

jal address // R[31] \leftarrow PC+8; PC \leftarrow address;

// Also, R-type jump

jr rs // PC \leftarrow R[rs]

Implementation of the MIPS ISA

- Hardware building blocks



Program Counter

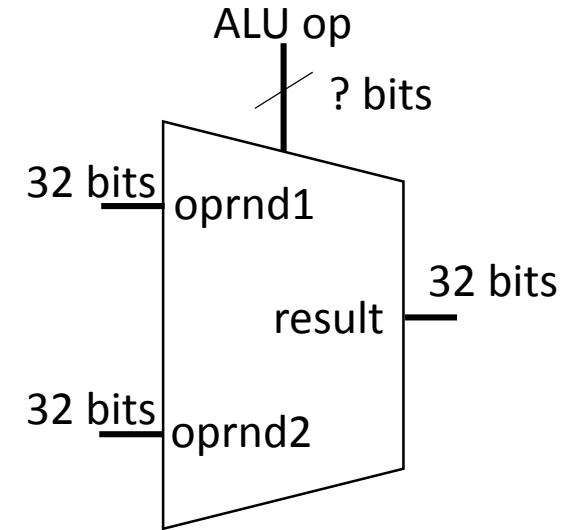
A 32-bit "register"
(can store information)

Register File

32 32-bit registers
(2 Read ports, 1 Write Port)

Set to '1' to
enable write

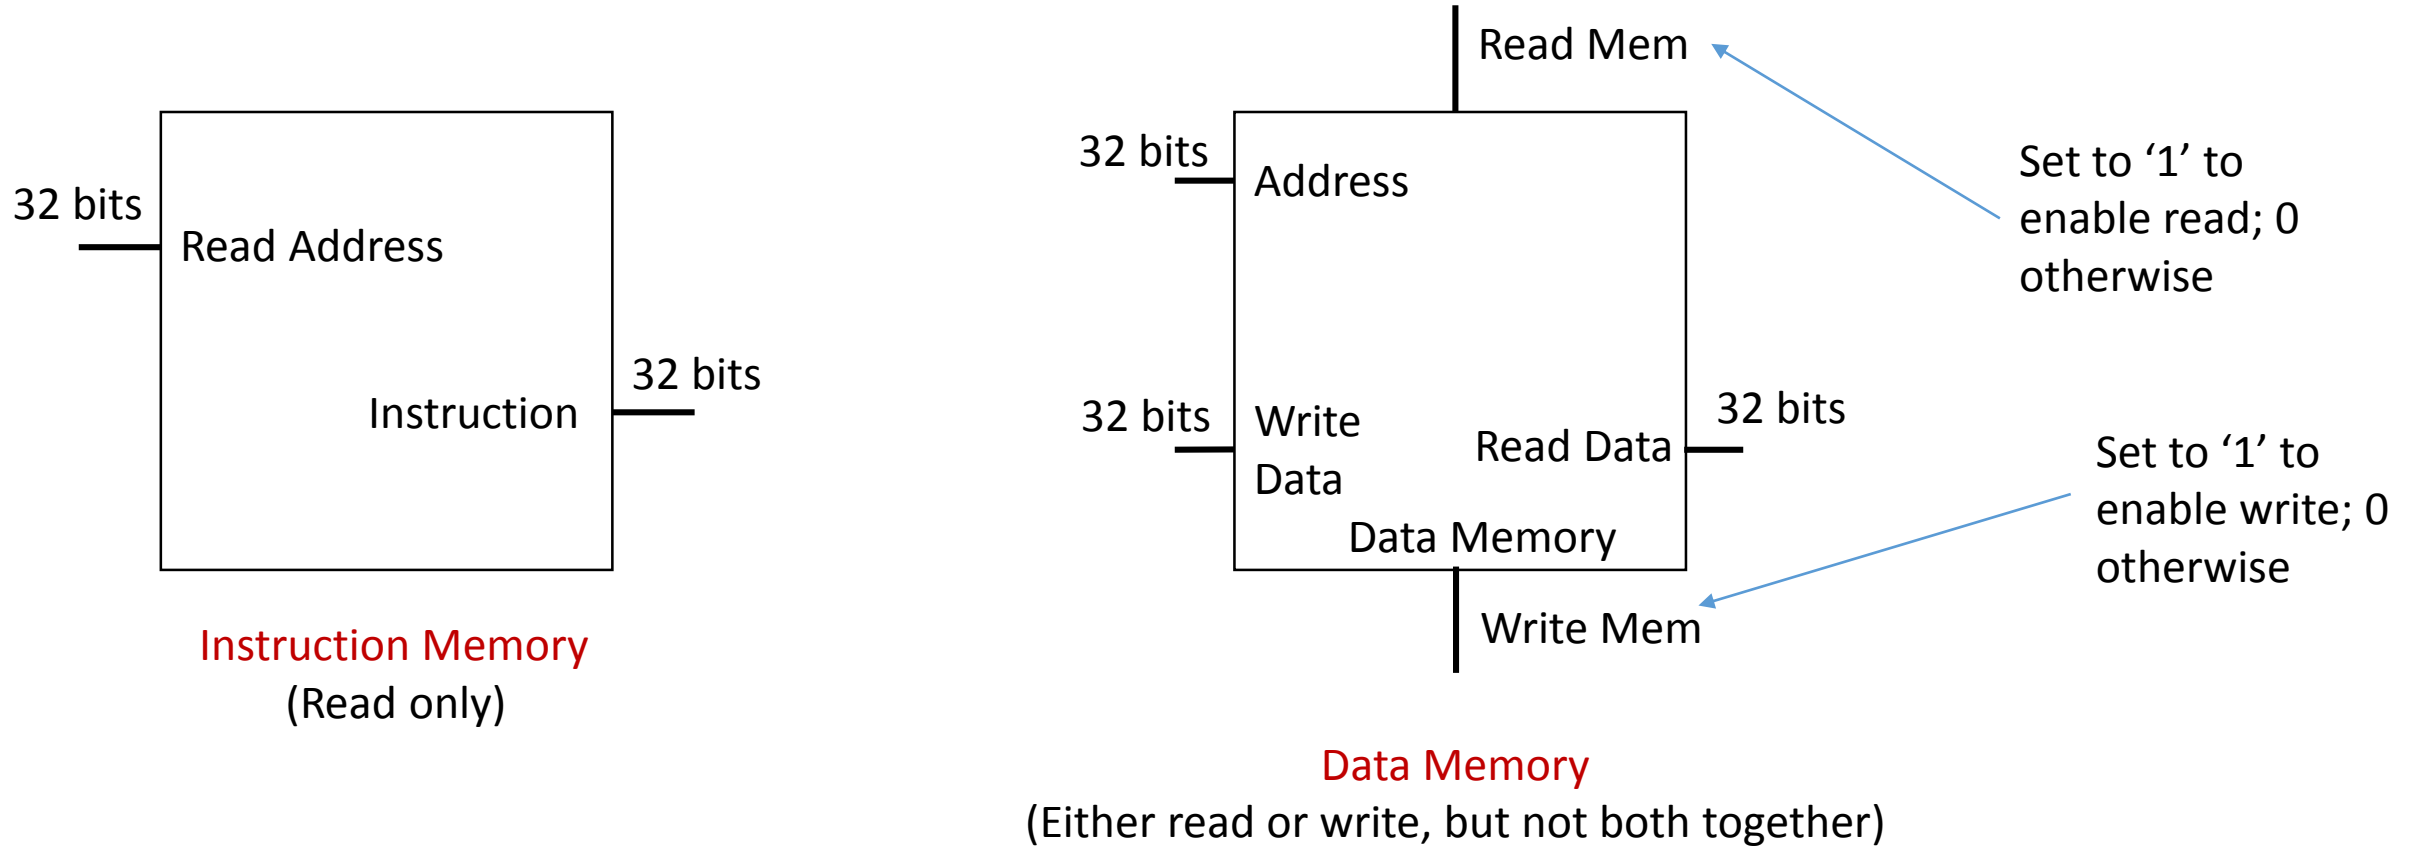
Arithmetic/Logic Unit (ALU)



Add: 20_{hex}
Addu: 20_{hex}
Sub: 23_{hex}
Subu: 24_{hex}
And: 25_{hex}
Or: 24_{hx}
Nor: 27_{hex}

Implementation of the MIPS ISA

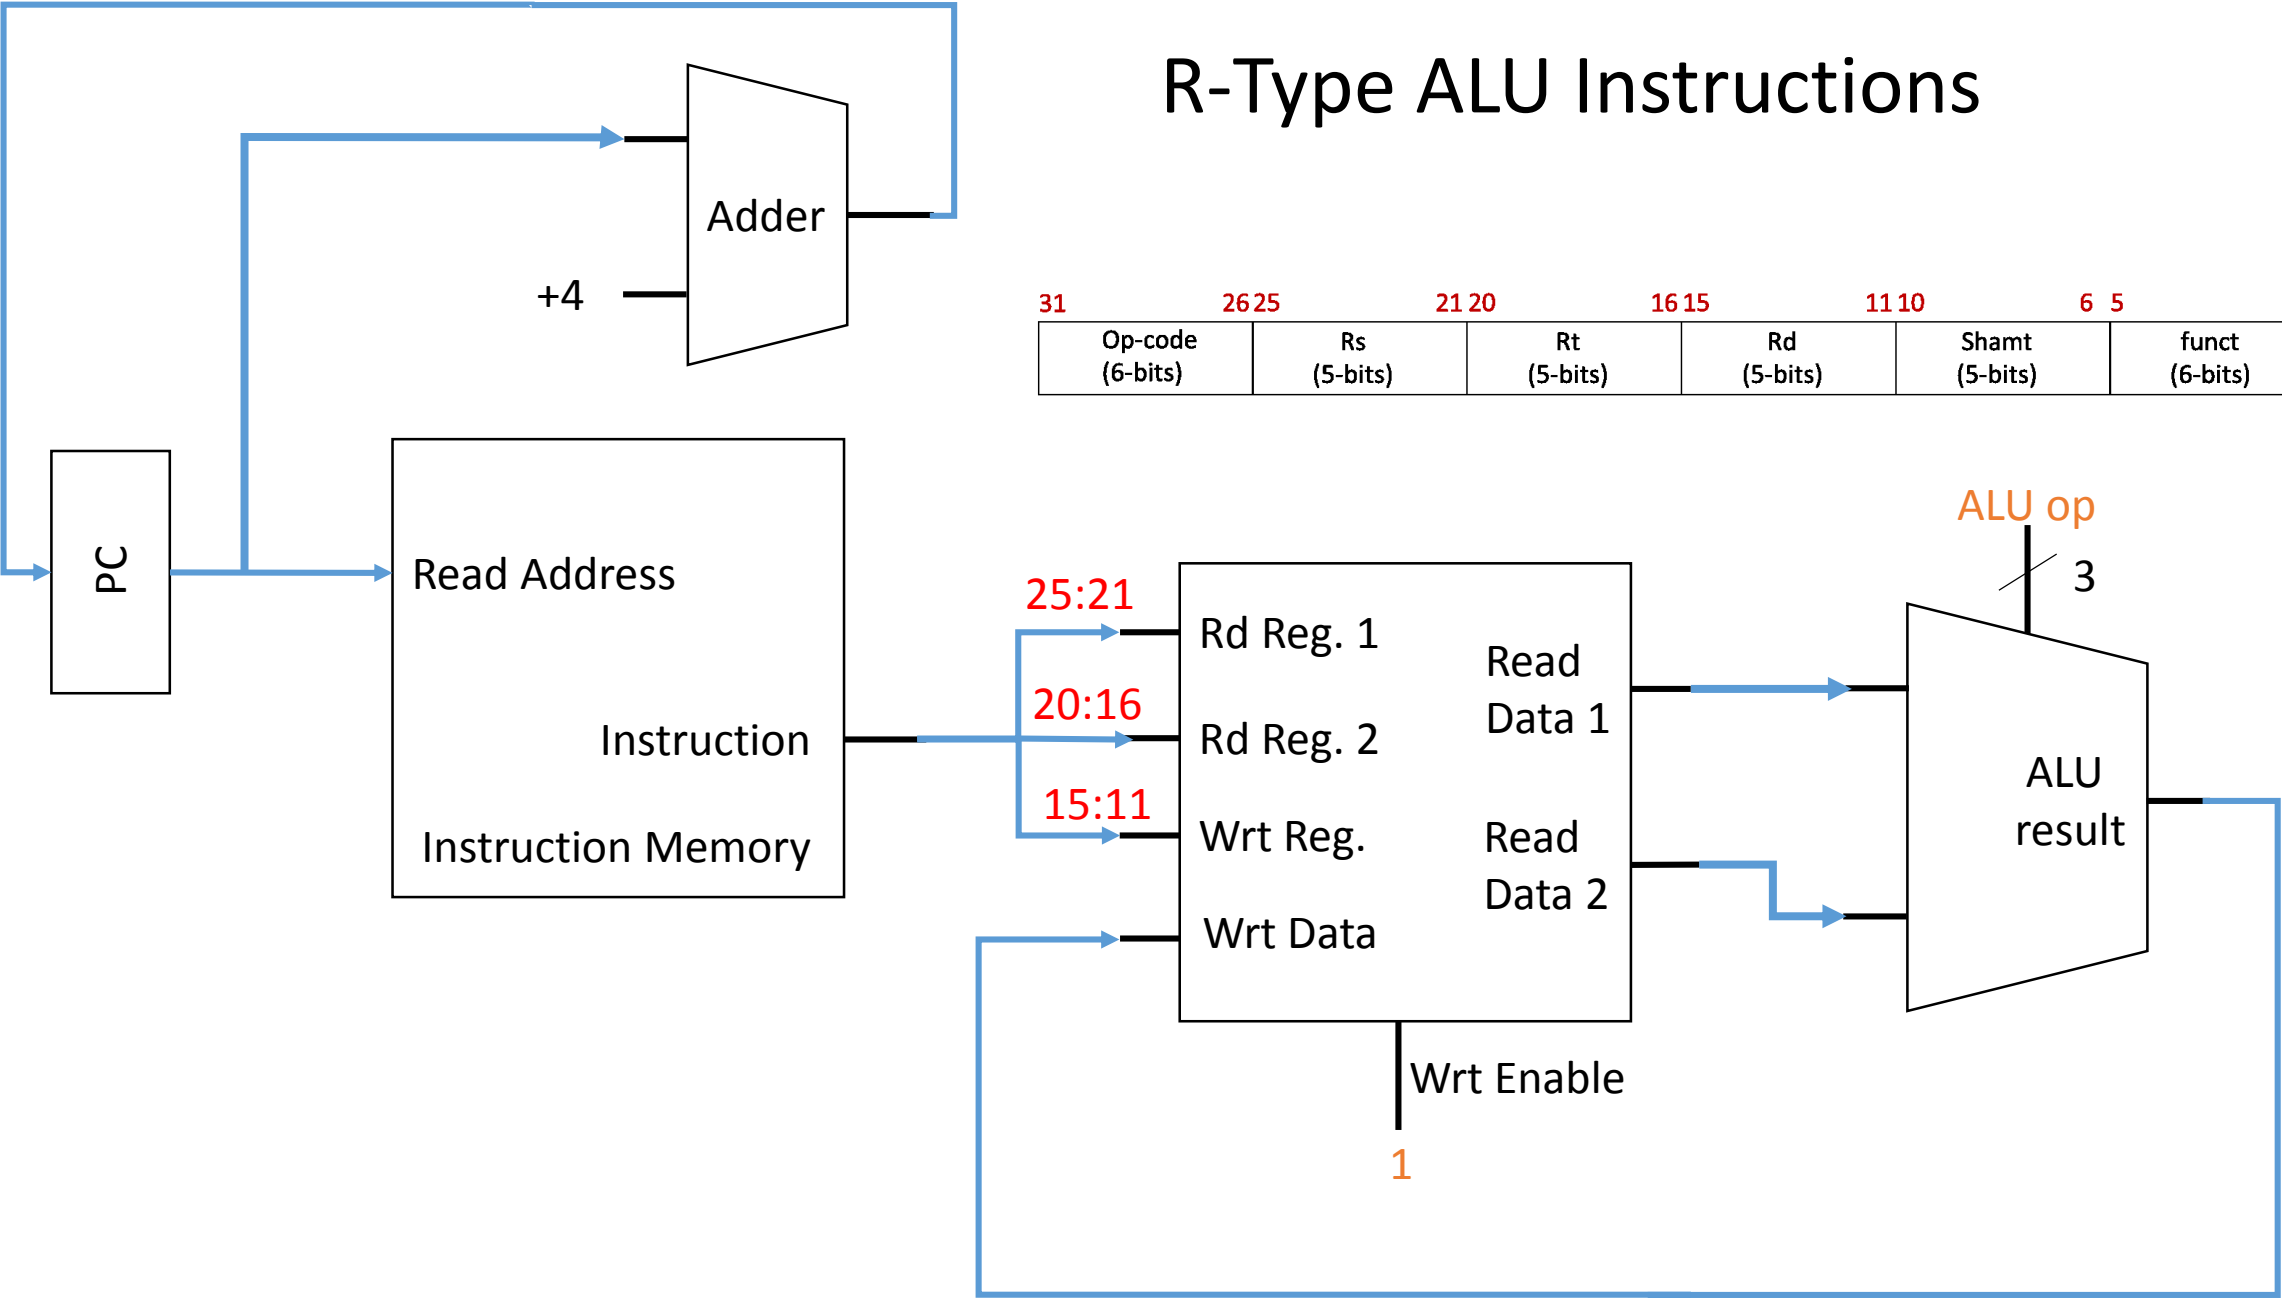
- Hardware building blocks



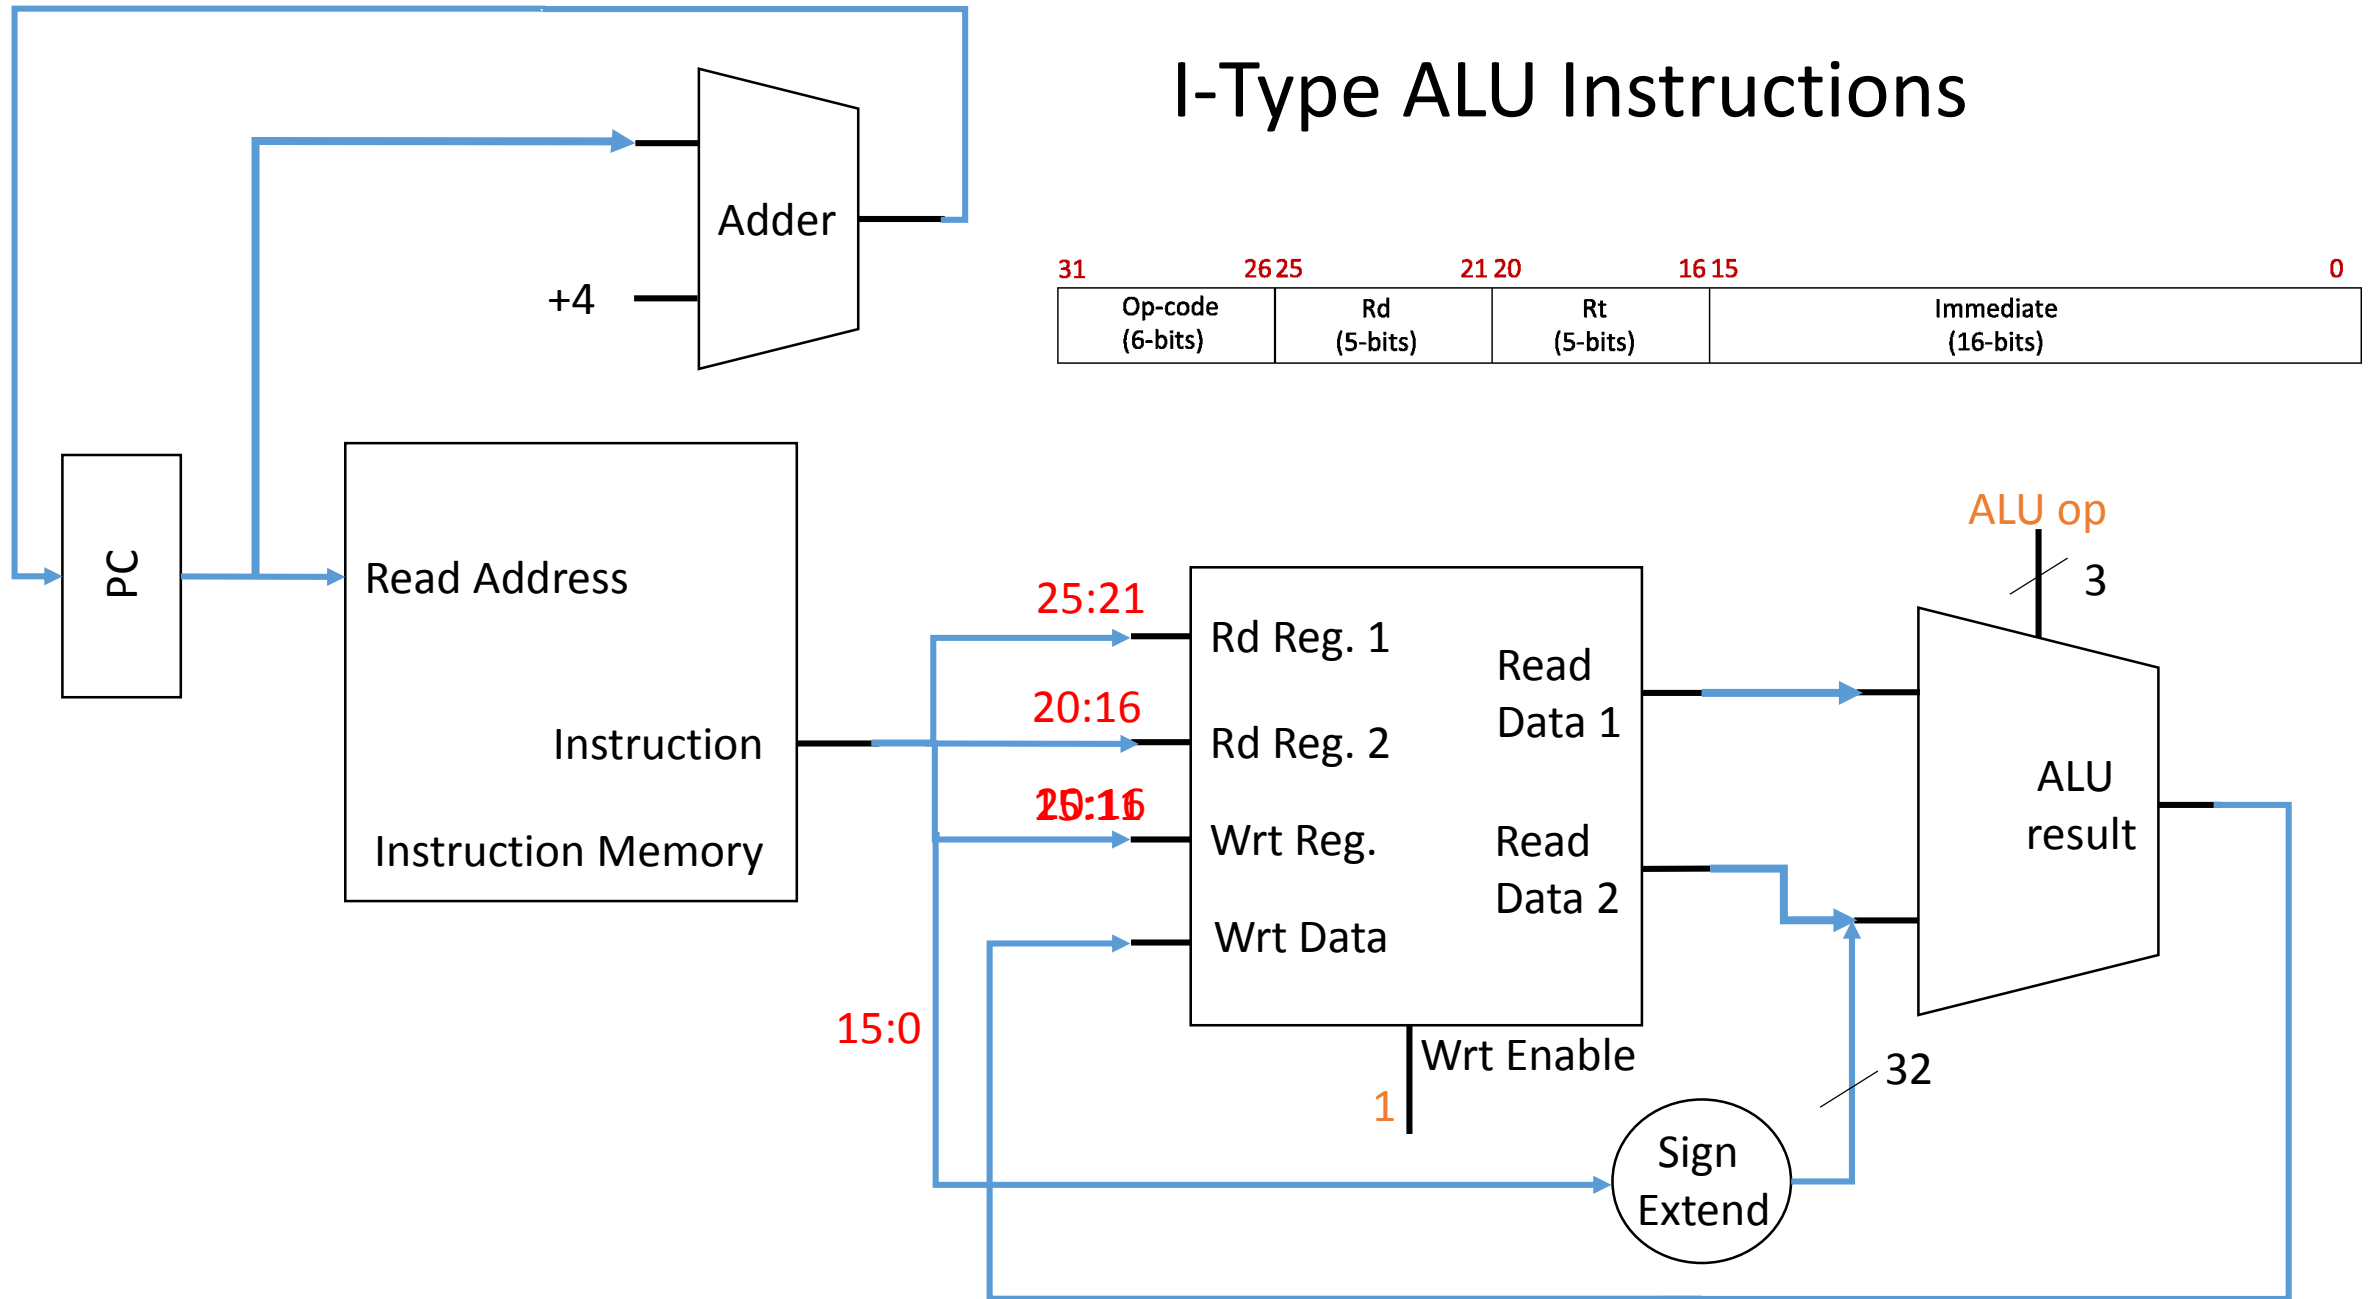
- Is this a von Neumann architecture?

R-Type ALU Instructions

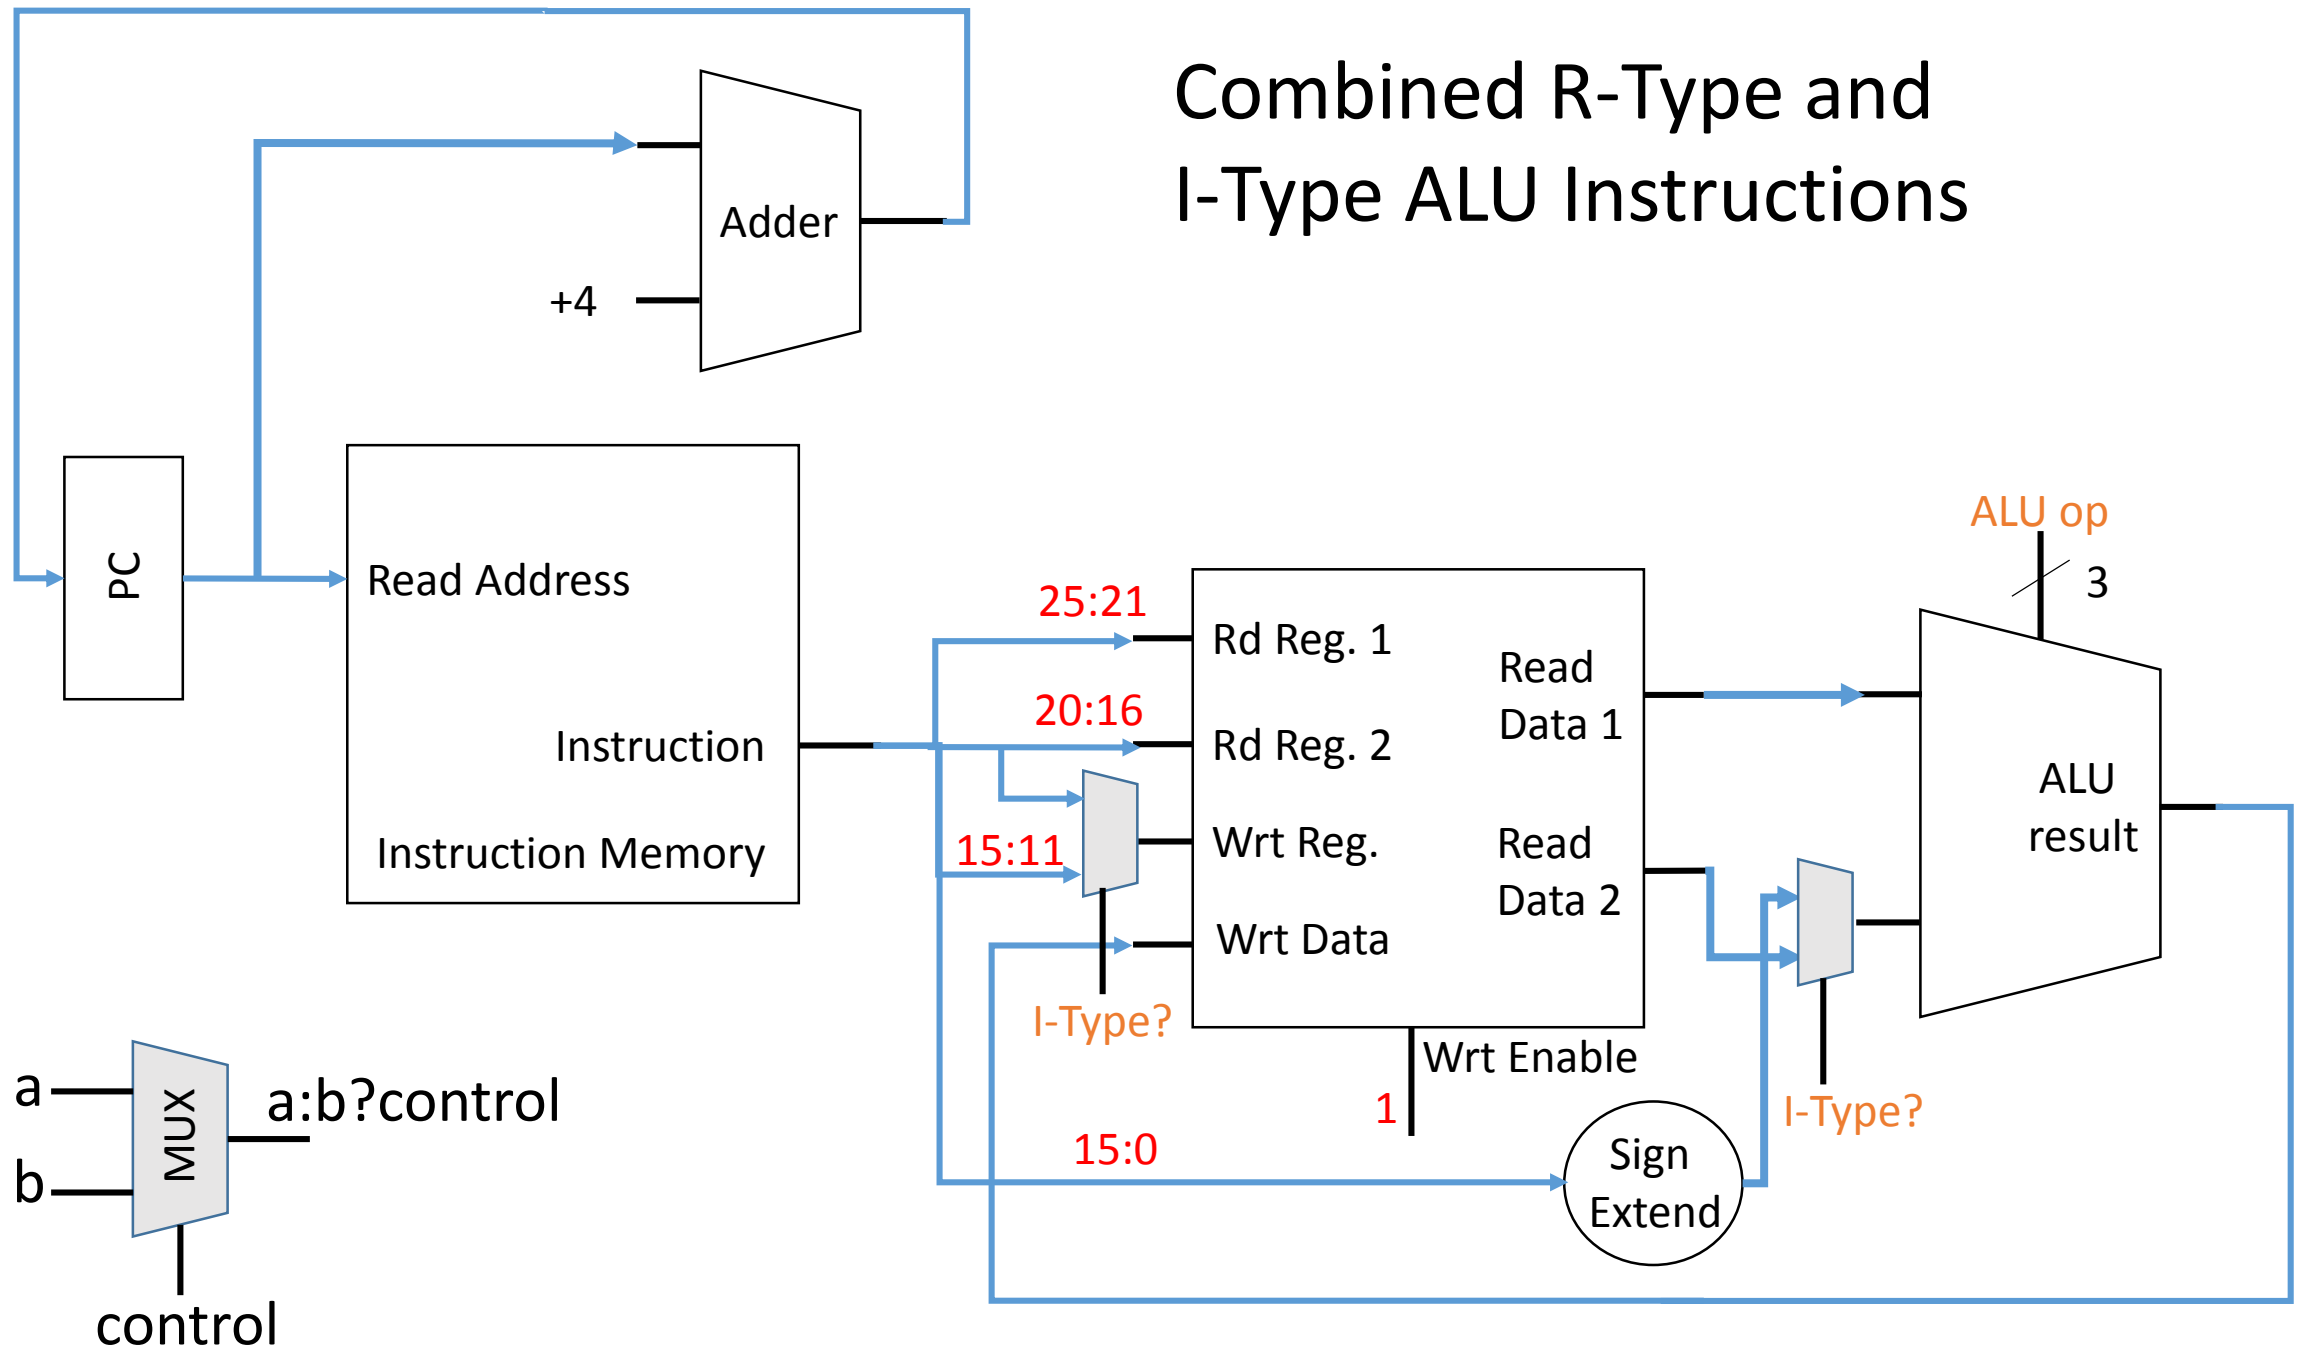
31	26 25	21 20	16 15	11 10	6 5	0
Op-code (6-bits)	Rs (5-bits)	Rt (5-bits)	Rd (5-bits)	Shamt (5-bits)	funct (6-bits)	



I-Type ALU Instructions

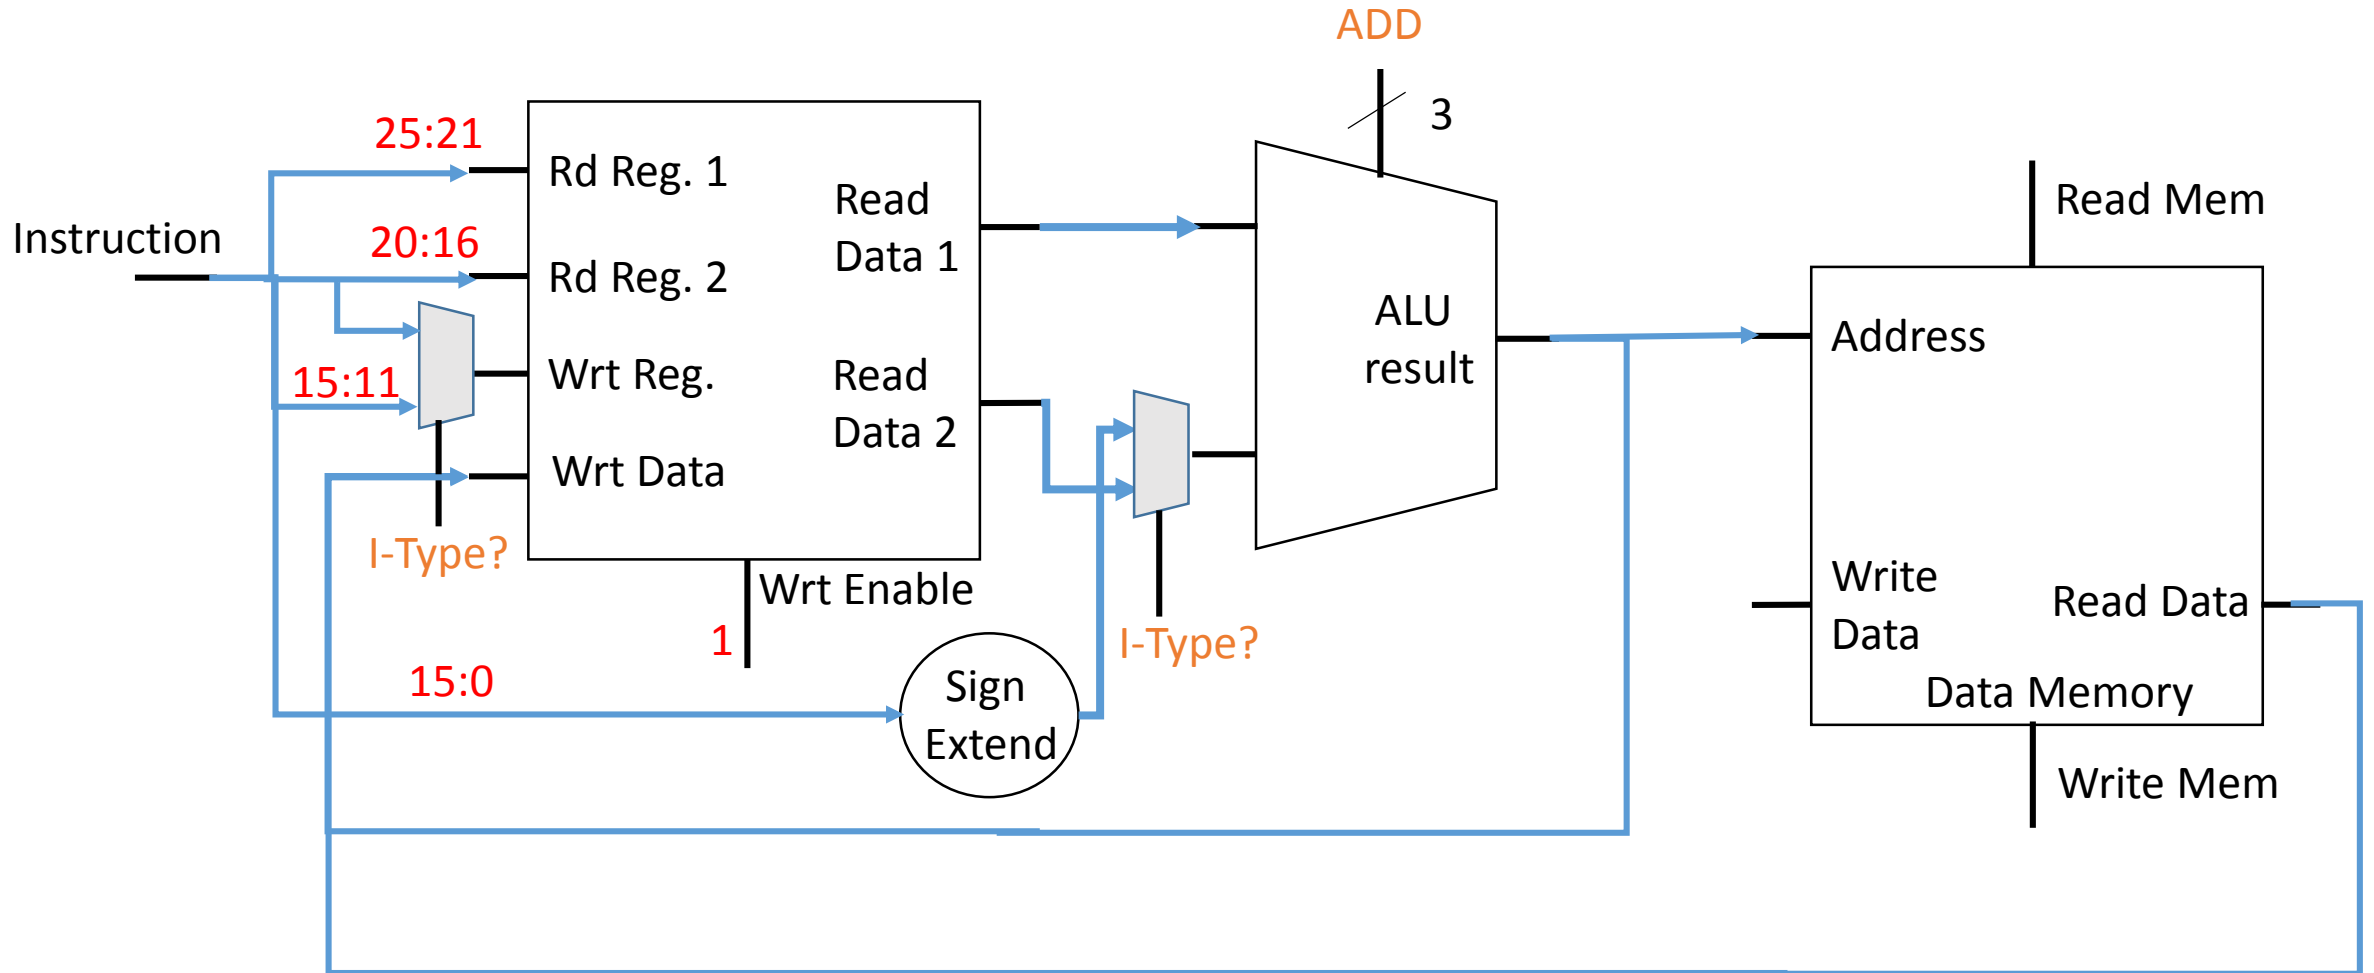


Combined R-Type and I-Type ALU Instructions



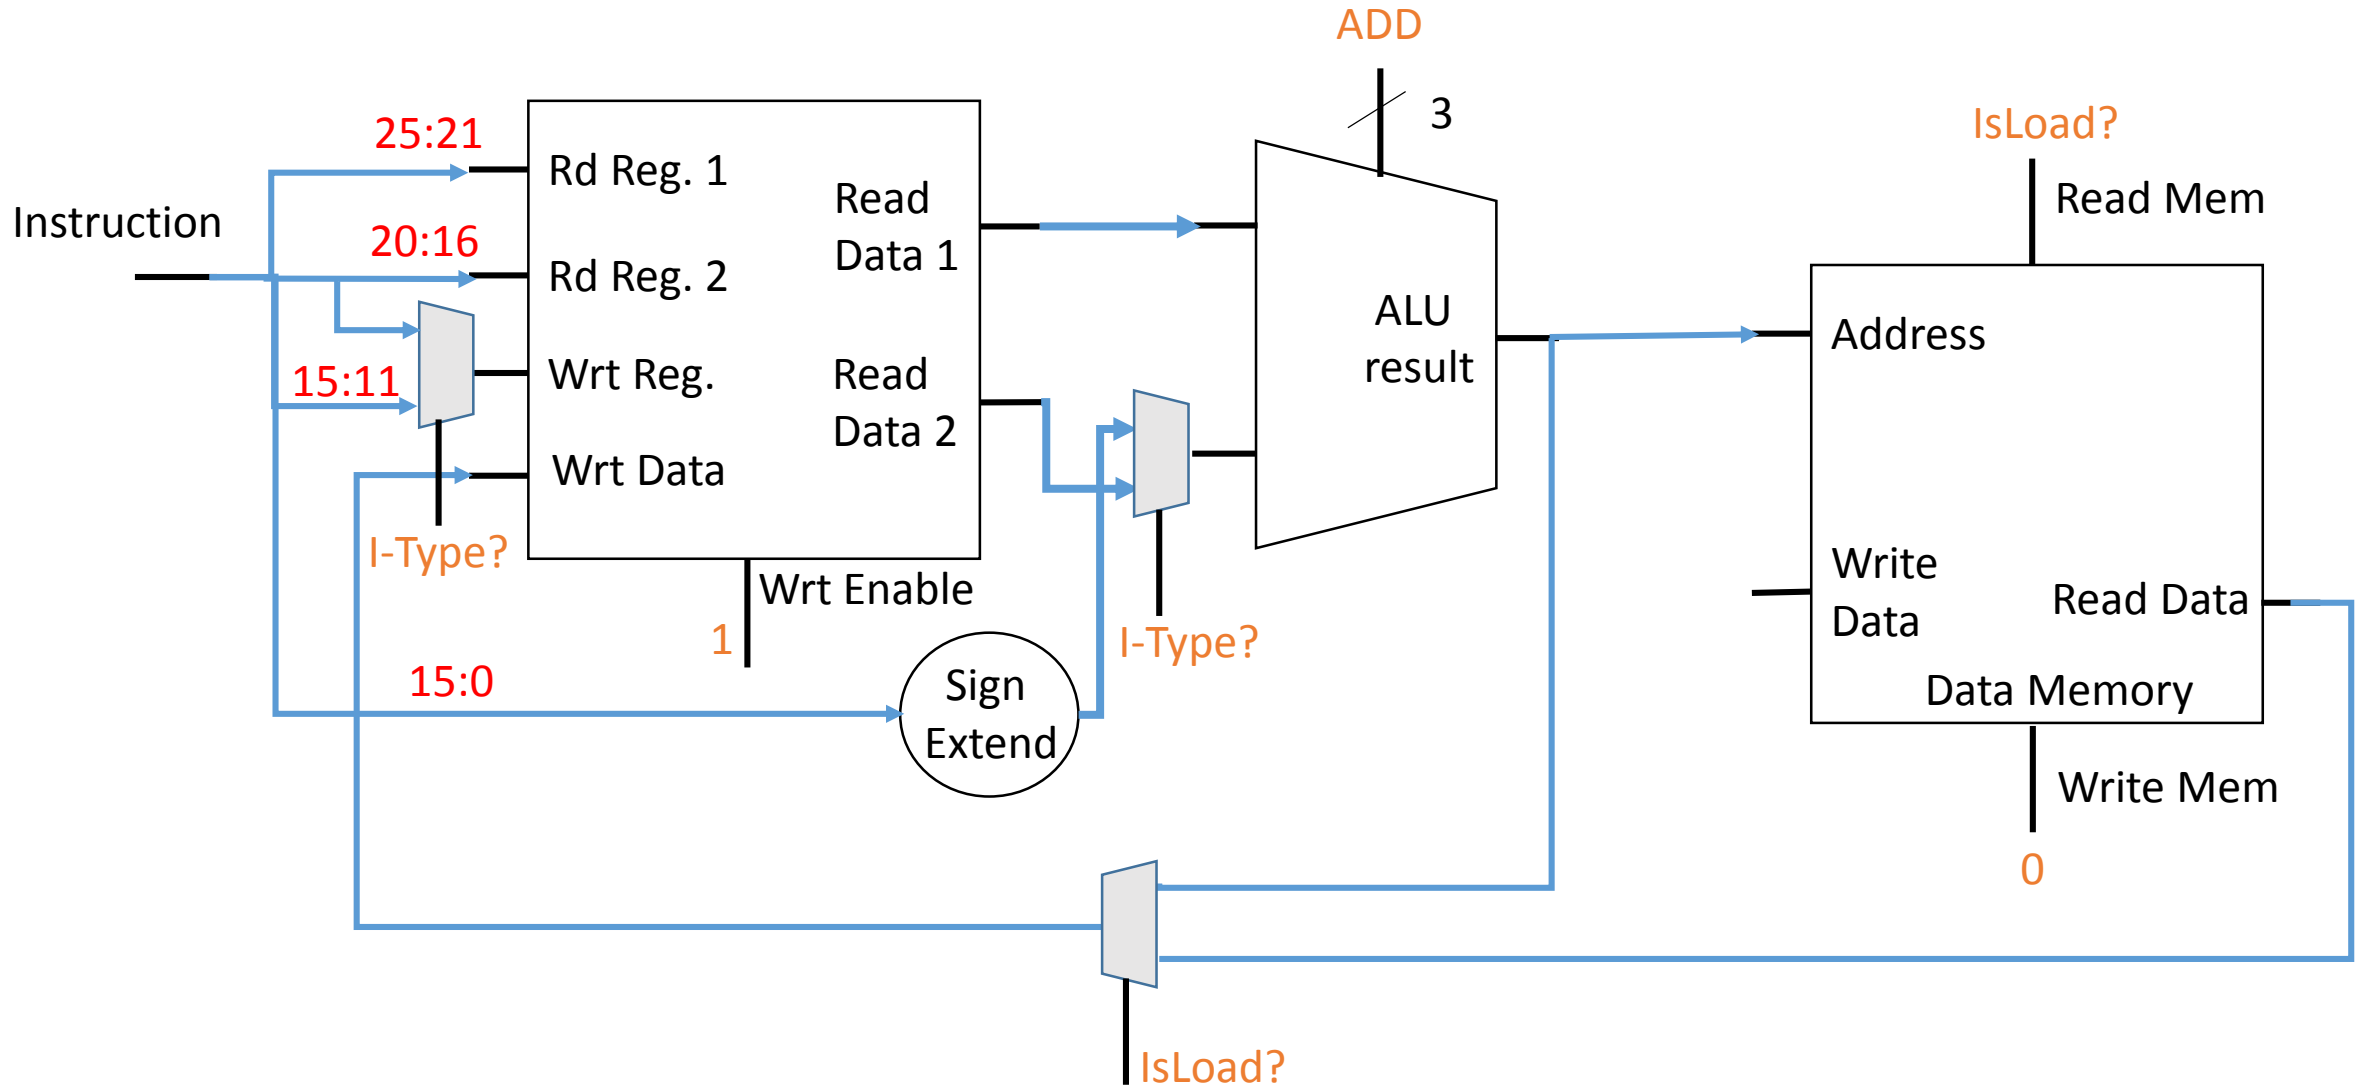
I-Type Load Instructions

`lw rs, rt, imm` $// R[rt] \leftarrow Mem[\{SignExtendimm\} + R[rs]]$ (“Displaced/based”)



I-Type Load Instructions

Instruction Fetch Part Not Shown



I-Type Store Instructions

Did we miss something?

